

Assignment-14-Decision_Tree - Company_data

```
In [79]: 1 # Importing Libraries
2 import pandas as pd
3 import numpy as np
4 import matplotlib.pyplot as plt
5 import seaborn as sns
6 %matplotlib inline
7 from sklearn.preprocessing import LabelEncoder # for encoding
8 from sklearn.model_selection import train_test_split # for train test splitting
9 from sklearn.tree import DecisionTreeClassifier # for decision tree object
10 from sklearn.metrics import classification_report, confusion_matrix # for checking testing results
11 from sklearn.tree import plot_tree # for visualizing tree
```

```
In [80]: 1 # Importing data
2 df = pd.read_csv('Company_Data.csv')
3 df.head()
```

```
Out[80]:
```

	Sales	CompPrice	Income	Advertising	Population	Price	ShelveLoc	Age	Education	Urban	US
0	9.50	138	73	11	276	120	Bad	42	17	Yes	Yes
1	11.22	111	48	16	260	83	Good	65	10	Yes	Yes
2	10.06	113	35	10	269	80	Medium	59	12	Yes	Yes
3	7.40	117	100	4	466	97	Medium	55	14	Yes	Yes
4	4.15	141	64	3	340	128	Bad	38	13	Yes	No

```
In [81]: 1 # Getting information of dataset
2 df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 400 entries, 0 to 399
Data columns (total 11 columns):
#   Column      Non-Null Count  Dtype
---  -
0   Sales       400 non-null    float64
1   CompPrice   400 non-null    int64
2   Income      400 non-null    int64
3   Advertising 400 non-null    int64
4   Population  400 non-null    int64
5   Price       400 non-null    int64
6   ShelveLoc   400 non-null    object
7   Age         400 non-null    int64
8   Education   400 non-null    int64
9   Urban       400 non-null    object
10  US          400 non-null    object
dtypes: float64(1), int64(7), object(3)
memory usage: 34.5+ KB
```

```
In [82]: 1 df.shape
```

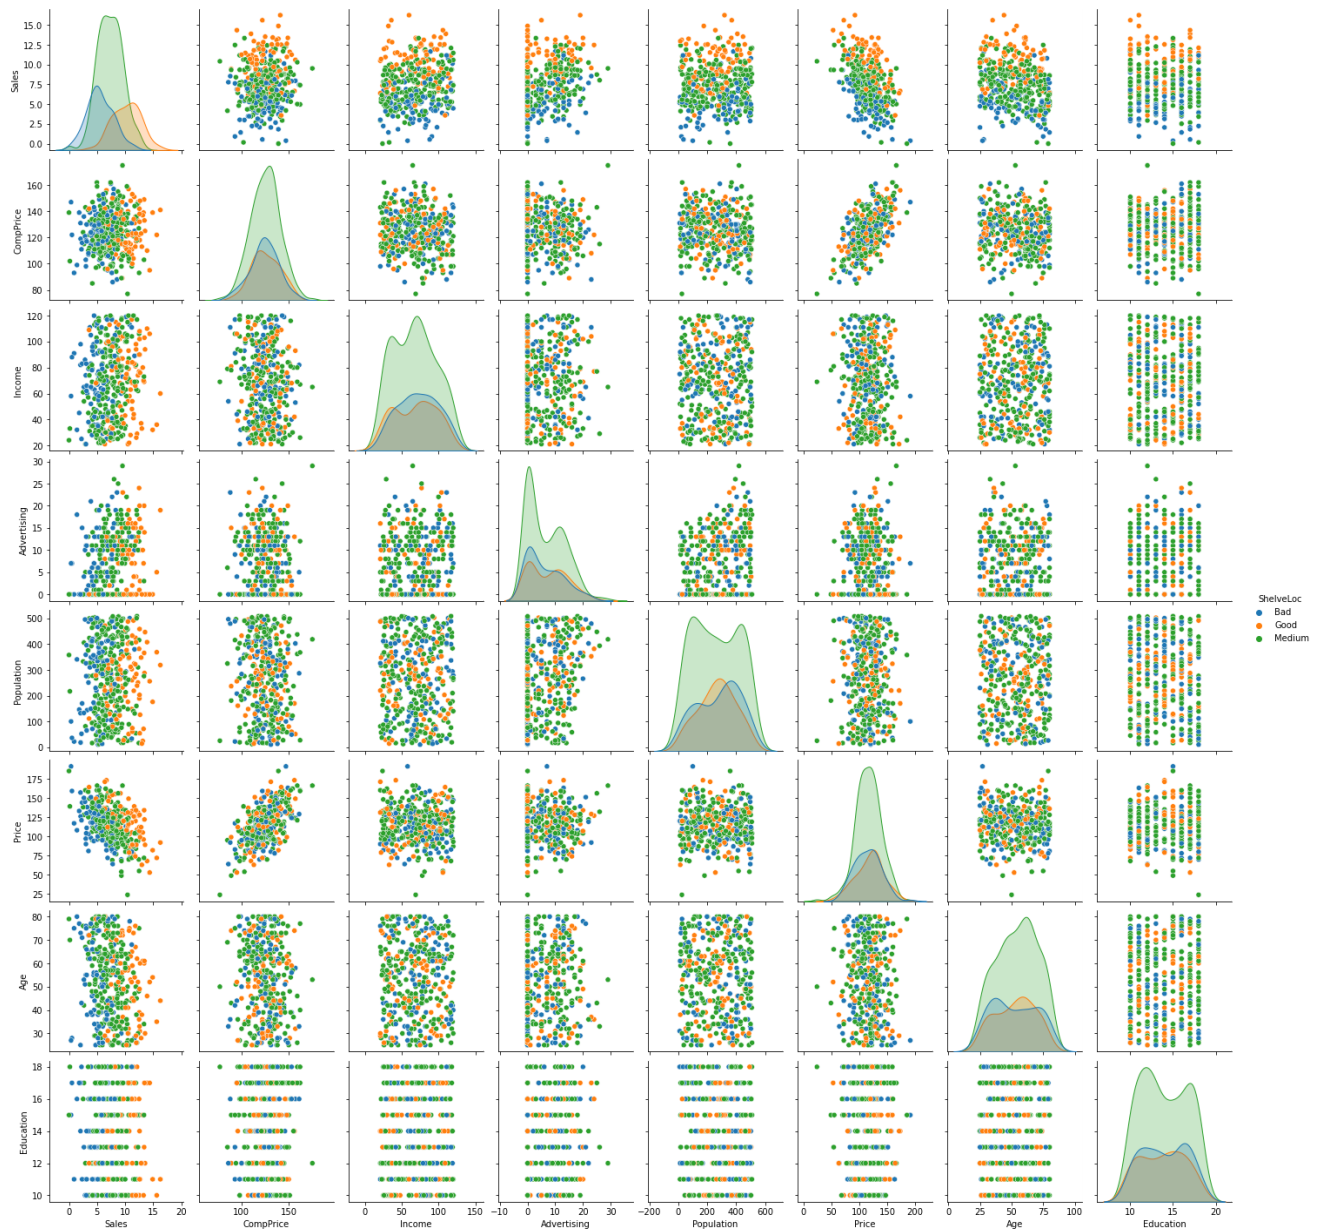
```
Out[82]: (400, 11)
```

```
In [83]: 1 df.isnull().any()
```

```
Out[83]: Sales      False
CompPrice  False
Income     False
Advertising False
Population False
Price      False
ShelveLoc  False
Age        False
Education  False
Urban      False
US         False
dtype: bool
```

```
In [84]: 1 # Lets plot pair plot to visualize the attributes all at once
        2 sns.pairplot(data=df, hue = 'ShelveLoc')
```

Out[84]: <seaborn.axisgrid.PairGrid at 0x229fd845d00>



```
In [85]: 1 # Creating dummy variables dropping first dummy variable
        2 df=pd.get_dummies(df,columns=['Urban','US'], drop_first=True)
```

In [86]:

1 df

Out[86]:

	Sales	CompPrice	Income	Advertising	Population	Price	ShelveLoc	Age	Education	Urban_Yes	US_Yes
0	9.50	138	73	11	276	120	Bad	42	17	1	1
1	11.22	111	48	16	260	83	Good	65	10	1	1
2	10.06	113	35	10	269	80	Medium	59	12	1	1
3	7.40	117	100	4	466	97	Medium	55	14	1	1
4	4.15	141	64	3	340	128	Bad	38	13	1	0
...
395	12.57	138	108	17	203	128	Good	33	14	1	1
396	6.14	139	23	3	37	120	Medium	55	11	0	1
397	7.41	162	26	12	368	159	Medium	40	18	1	1
398	5.94	100	79	7	284	95	Bad	50	12	1	1
399	9.71	134	37	0	27	120	Good	49	16	1	1

400 rows × 11 columns

In [87]:

1 df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 400 entries, 0 to 399
Data columns (total 11 columns):
#   Column          Non-Null Count  Dtype
---  ---
0   Sales            400 non-null    float64
1   CompPrice        400 non-null    int64
2   Income           400 non-null    int64
3   Advertising       400 non-null    int64
4   Population        400 non-null    int64
5   Price            400 non-null    int64
6   ShelveLoc        400 non-null    object
7   Age              400 non-null    int64
8   Education         400 non-null    int64
9   Urban_Yes        400 non-null    uint8
10  US_Yes           400 non-null    uint8
dtypes: float64(1), int64(7), object(1), uint8(2)
memory usage: 29.0+ KB
```

In [88]:

```
1 from sklearn.metrics import f1_score
2 from sklearn.model_selection import train_test_split
```

In [89]:

1 df['ShelveLoc']=df['ShelveLoc'].map({'Good':1,'Medium':2,'Bad':3})

In [90]:

1 df.head()

Out[90]:

	Sales	CompPrice	Income	Advertising	Population	Price	ShelveLoc	Age	Education	Urban_Yes	US_Yes
0	9.50	138	73	11	276	120	3	42	17	1	1
1	11.22	111	48	16	260	83	1	65	10	1	1
2	10.06	113	35	10	269	80	2	59	12	1	1
3	7.40	117	100	4	466	97	2	55	14	1	1
4	4.15	141	64	3	340	128	3	38	13	1	0

In [91]:

```
1 x=df.iloc[:,0:6]
2 y=df['ShelveLoc']
```

In [92]:

1 x

Out[92]:

	Sales	CompPrice	Income	Advertising	Population	Price
0	9.50	138	73	11	276	120
1	11.22	111	48	16	260	83
2	10.06	113	35	10	269	80
3	7.40	117	100	4	466	97
4	4.15	141	64	3	340	128
...
395	12.57	138	108	17	203	128
396	6.14	139	23	3	37	120
397	7.41	162	26	12	368	159
398	5.94	100	79	7	284	95
399	9.71	134	37	0	27	120

400 rows × 6 columns

In [93]:

1 y

Out[93]:

```

0    3
1    1
2    2
3    2
4    3
..
395   1
396   2
397   2
398   3
399   1
Name: ShelveLoc, Length: 400, dtype: int64

```

In [94]:

1 df['ShelveLoc'].unique()

Out[94]:

array([3, 1, 2], dtype=int64)

In [95]:

1 df.ShelveLoc.value_counts()

Out[95]:

```

2    219
3     96
1     85
Name: ShelveLoc, dtype: int64

```

In [96]:

```

1 colnames = list(df.columns)
2 colnames

```

Out[96]:

```

['Sales',
 'CompPrice',
 'Income',
 'Advertising',
 'Population',
 'Price',
 'ShelveLoc',
 'Age',
 'Education',
 'Urban_Yes',
 'US_Yes']

```

In [97]:

```

1 # Splitting data into training and testinh dataset
2 x_train, x_test, y_train, y_test = train_test_split(x,y, test_size=0.2,random_state=40)

```

Building Decision Tree Classifier using Entropy Criteria.

In [98]:

```

1 model = DecisionTreeClassifier(criterion = 'entropy',max_depth=3)
2 model.fit(x_train,y_train)

```


Out[98]:

DecisionTreeClassifier(criterion='entropy', max_depth=3)

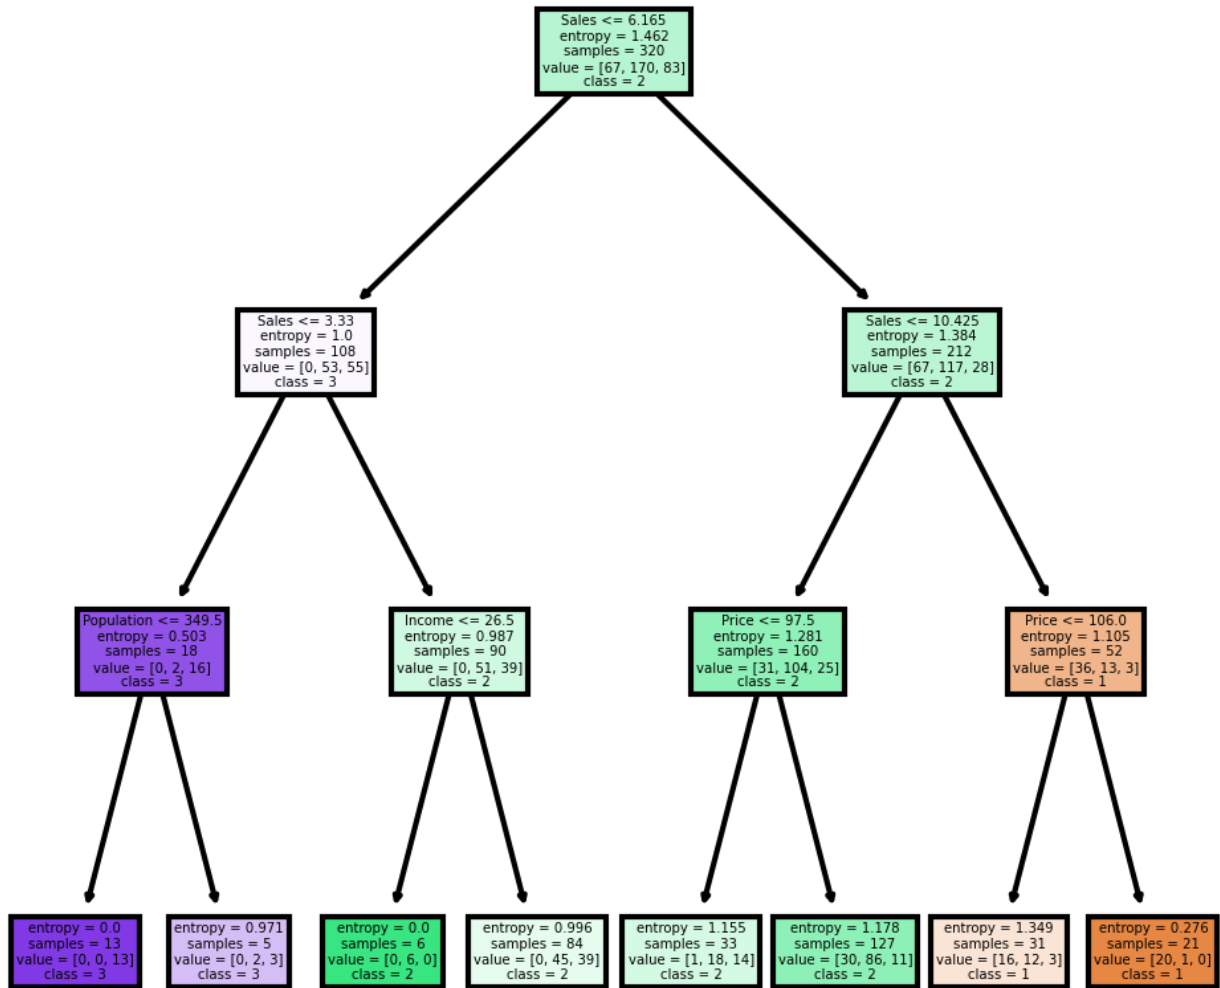
In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.
On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
In [99]: 1 from sklearn import tree

In [100]: 1 # Plot the decision tree
          2 tree.plot_tree(model);
```



```
In [101]: 1 fn=['Sales','CompPrice','Income','Advertising','Population','Price']
2 cn=['1', '2', '3']
3 fig, axes = plt.subplots(nrows = 1,ncols = 1,figsize = (4,4), dpi=300)
4 tree.plot_tree(model,
5                 feature_names = fn,
6                 class_names = cn,
7                 filled = True);
```



```
In [102]: 1 # Predicting on test data
2 preds = model.predict(x_test) # predicting on test dataset
3 pd.Series(preds).value_counts() # getting the count of each category
```

```
Out[102]: 2 63
1 13
3 4
dtype: int64
```

```
In [103]: 1 preds
```

```
Out[103]: array([2, 2, 2, 2, 2, 2, 2, 1, 2, 1, 2, 2, 2, 2, 2, 1, 2, 2, 2, 2, 2, 1,
2, 2, 2, 2, 2, 2, 1, 1, 1, 2, 1, 2, 2, 1, 2, 2, 2, 2, 2, 2, 2, 3, 2,
2, 2, 2, 2, 2, 2, 3, 2, 2, 2, 2, 1, 1, 2, 2, 3, 2, 2, 1, 2, 2, 2,
2, 2, 2, 1, 2, 2, 2, 2, 2, 2, 2, 3, 2, 2, 2], dtype=int64)
```

```
In [104]: 1 pd.crosstab(y_test,preds) # getting the 2 way table to understand the correct and wrong predictions
```

```
Out[104]:      col_0  1   2   3
      ShelfLoc
          1   8  10   0
          2   5  41   3
          3   0  12   1
```

```
In [105]: 1 # Accuracy
          2 np.mean(preds==y_test)
```

```
Out[105]: 0.625
```

Building Decision Tree Classifier (CART) using Gini Criteria

```
In [106]: 1 from sklearn.tree import DecisionTreeClassifier
          2 model_gini = DecisionTreeClassifier(criterion='gini', max_depth=3)
          3 model_gini.fit(x_train, y_train)
```

```
Out[106]: DecisionTreeClassifier(max_depth=3)
```

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.
On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
In [107]: 1 # Prediction and computing the accuracy
          2 pred = model.predict(x_test)
          3 np.mean(preds==y_test)
```

```
Out[107]: 0.625
```

Decision Tree Regression Example

```
In [108]: 1 # Decision Tree Regression
          2 from sklearn.tree import DecisionTreeRegressor
```

```
In [109]: 1 array = df.values
          2 X = array[:,0:3]
          3 y = array[:,3]
```

```
In [110]: 1 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.33, random_state=1)
```

```
In [111]: 1 model = DecisionTreeRegressor()
          2 model.fit(X_train, y_train)
```

```
Out[111]: DecisionTreeRegressor()
```

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.
On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
In [112]: 1 # Find the accuracy
          2 model.score(X_test,y_test)
```

```
Out[112]: -1.2646168892692322
```

```
In [ ]: 1
```

```
In [ ]: 1
```