

Assignment-06-Logistic Regression

In [5]:

```
1 # Import the Libraries
2 import pandas as pd
3 import numpy as np
4 import matplotlib.pyplot as plt
5 import seaborn as sns
6 from sklearn.linear_model import LogisticRegression
7 from sklearn.metrics import confusion_matrix
8 from sklearn.metrics import roc_curve
9 from sklearn.metrics import roc_auc_score
```

In [6]:

```
1 # Importing the dataset
2 bank=pd.read_csv('bank-full.csv', sep=';')
3 bank
```

Out[6]:

	age	job	marital	education	default	balance	housing	loan	contact	day	month	duration	campaign	pdays	previous	poutcome	y	
	0	58	management	married	tertiary	no	2143	yes	no	unknown	5	may	261	1	-1	0	unknown	no
	1	44	technician	single	secondary	no	29	yes	no	unknown	5	may	151	1	-1	0	unknown	no
	2	33	entrepreneur	married	secondary	no	2	yes	yes	unknown	5	may	76	1	-1	0	unknown	no
	3	47	blue-collar	married	unknown	no	1506	yes	no	unknown	5	may	92	1	-1	0	unknown	no
	4	33	unknown	single	unknown	no	1	no	no	unknown	5	may	198	1	-1	0	unknown	no

	45206	51	technician	married	tertiary	no	825	no	no	cellular	17	nov	977	3	-1	0	unknown	yes
	45207	71	retired	divorced	primary	no	1729	no	no	cellular	17	nov	456	2	-1	0	unknown	yes
	45208	72	retired	married	secondary	no	5715	no	no	cellular	17	nov	1127	5	184	3	success	yes
	45209	57	blue-collar	married	secondary	no	668	no	no	telephone	17	nov	508	4	-1	0	unknown	no
	45210	37	entrepreneur	married	secondary	no	2971	no	no	cellular	17	nov	361	2	188	11	other	no

45211 rows × 17 columns

EDA

In [7]:

```
1 bank.info()
```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 45211 entries, 0 to 45210
Data columns (total 17 columns):
Column Non-Null Count Dtype
--- ---
0 age 45211 non-null int64
1 job 45211 non-null object
2 marital 45211 non-null object
3 education 45211 non-null object
4 default 45211 non-null object
5 balance 45211 non-null int64
6 housing 45211 non-null object
7 loan 45211 non-null object
8 contact 45211 non-null object
9 day 45211 non-null int64
10 month 45211 non-null object
11 duration 45211 non-null int64
12 campaign 45211 non-null int64
13 pdays 45211 non-null int64
14 previous 45211 non-null int64
15 poutcome 45211 non-null object
16 y 45211 non-null object
dtypes: int64(7), object(10)
memory usage: 5.9+ MB

```
In [8]: 1 # One-Hot Encoding of Categorical variables
2 data1=pd.get_dummies(bank,columns=['job','marital','education','contact','outcome'])
3 data1
```

Out[8]:

	age	default	balance	housing	loan	day	month	duration	campaign	pdays	...	education_secondary	education_tertiary	education_unknown	contact
0	58	no	2143	yes	no	5	may	261	1	-1	...	0	1	0	
1	44	no	29	yes	no	5	may	151	1	-1	...	1	0	0	
2	33	no	2	yes	yes	5	may	76	1	-1	...	1	0	0	
3	47	no	1506	yes	no	5	may	92	1	-1	...	0	0	1	
4	33	no	1	no	no	5	may	198	1	-1	...	0	0	1	
...
45206	51	no	825	no	no	17	nov	977	3	-1	...	0	1	0	
45207	71	no	1729	no	no	17	nov	456	2	-1	...	0	0	0	
45208	72	no	5715	no	no	17	nov	1127	5	184	...	1	0	0	
45209	57	no	668	no	no	17	nov	508	4	-1	...	1	0	0	
45210	37	no	2971	no	no	17	nov	361	2	188	...	1	0	0	

45211 rows × 38 columns

```
In [9]: 1 # To see all columns
2 pd.set_option("display.max.columns",None)
3 data1
```

Out[9]:

	age	default	balance	housing	loan	day	month	duration	campaign	pdays	previous	y	job_admin.	job_blue-collar	job_entrepreneur	job_housemaid	j
0	58	no	2143	yes	no	5	may	261	1	-1	0	no	0	0	0	0	
1	44	no	29	yes	no	5	may	151	1	-1	0	no	0	0	0	0	
2	33	no	2	yes	yes	5	may	76	1	-1	0	no	0	0	1	0	
3	47	no	1506	yes	no	5	may	92	1	-1	0	no	0	1	0	0	
4	33	no	1	no	no	5	may	198	1	-1	0	no	0	0	0	0	
...
45206	51	no	825	no	no	17	nov	977	3	-1	0	yes	0	0	0	0	
45207	71	no	1729	no	no	17	nov	456	2	-1	0	yes	0	0	0	0	
45208	72	no	5715	no	no	17	nov	1127	5	184	3	yes	0	0	0	0	
45209	57	no	668	no	no	17	nov	508	4	-1	0	no	0	1	0	0	
45210	37	no	2971	no	no	17	nov	361	2	188	11	no	0	0	1	0	

45211 rows × 38 columns

In [10]:

1

data1.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 45211 entries, 0 to 45210
Data columns (total 38 columns):
Column Non-Null Count Dtype
--- -
0 age 45211 non-null int64
1 default 45211 non-null object
2 balance 45211 non-null int64
3 housing 45211 non-null object
4 loan 45211 non-null object
5 day 45211 non-null int64
6 month 45211 non-null object
7 duration 45211 non-null int64
8 campaign 45211 non-null int64
9 pdays 45211 non-null int64
10 previous 45211 non-null int64
11 y 45211 non-null object
12 job_admin. 45211 non-null uint8
13 job_blue-collar 45211 non-null uint8
14 job_entrepreneur 45211 non-null uint8
15 job_housemaid 45211 non-null uint8
16 job_management 45211 non-null uint8
17 job_retired 45211 non-null uint8
18 job_self-employed 45211 non-null uint8
19 job_services 45211 non-null uint8
20 job_student 45211 non-null uint8
21 job_technician 45211 non-null uint8
22 job_unemployed 45211 non-null uint8
23 job_unknown 45211 non-null uint8
24 marital_divorced 45211 non-null uint8
25 marital_married 45211 non-null uint8
26 marital_single 45211 non-null uint8
27 education_primary 45211 non-null uint8
28 education_secondary 45211 non-null uint8
29 education_tertiary 45211 non-null uint8
30 education_unknown 45211 non-null uint8
31 contact_cellular 45211 non-null uint8
32 contact_telephone 45211 non-null uint8
33 contact_unknown 45211 non-null uint8
34 poutcome_failure 45211 non-null uint8
35 poutcome_other 45211 non-null uint8
36 poutcome_success 45211 non-null uint8
37 poutcome_unknown 45211 non-null uint8
dtypes: int64(7), object(5), uint8(26)
memory usage: 5.3+ MB

In [11]:

1 #Custom Binary Encoding of Binary o/p variables
2 data1['default'] = np.where(data1['default'].str.contains("yes"), 1, 0)
3 data1['housing'] = np.where(data1['housing'].str.contains("yes"), 1, 0)
4 data1['loan'] = np.where(data1['loan'].str.contains("yes"), 1, 0)
5 data1['y'] = np.where(data1['y'].str.contains("yes"), 1, 0)
6 data1

Out[11]:

	age	default	balance	housing	loan	day	month	duration	campaign	pdays	previous	y	job_admin.	job_blue-collar	job_entrepreneur	job_housemaid	jo
0	58	0	2143	1	0	5	may	261	1	-1	0	0	0	0	0	0	0
1	44	0	29	1	0	5	may	151	1	-1	0	0	0	0	0	0	0
2	33	0	2	1	1	5	may	76	1	-1	0	0	0	0	1	0	0
3	47	0	1506	1	0	5	may	92	1	-1	0	0	0	1	0	0	0
4	33	0	1	0	0	5	may	198	1	-1	0	0	0	0	0	0	0
...
45206	51	0	825	0	0	17	nov	977	3	-1	0	1	0	0	0	0	0
45207	71	0	1729	0	0	17	nov	456	2	-1	0	1	0	0	0	0	0
45208	72	0	5715	0	0	17	nov	1127	5	184	3	1	0	0	0	0	0
45209	57	0	668	0	0	17	nov	508	4	-1	0	0	0	1	0	0	0
45210	37	0	2971	0	0	17	nov	361	2	188	11	0	0	0	1	0	0

45211 rows × 38 columns

localhost:8888/notebooks/Assignment-06-Logistic Regression.ipynb

3/7

```
In [13]: 1 import warnings
2 warnings.filterwarnings('ignore')
3 from sklearn.preprocessing import LabelEncoder
4 for column in bank.columns:
5     if bank[column].dtype == np.number:
6         continue
7     bank[column] = LabelEncoder().fit_transform(bank[column])
8 bank.head()
```

Out[13]:

	age	job	marital	education	default	balance	housing	loan	contact	day	month	duration	campaign	pdays	previous	poutcome	y
0	40	4	1	2	0	3036	1	0	2	4	8	261	0	0	0	3	0
1	26	9	2	1	0	945	1	0	2	4	8	151	0	0	0	3	0
2	15	2	1	1	0	918	1	1	2	4	8	76	0	0	0	3	0
3	29	1	1	3	0	2420	1	0	2	4	8	92	0	0	0	3	0
4	15	11	2	3	0	917	0	0	2	4	8	198	0	0	0	3	0

```
In [14]: 1 # Find and Replace Encoding for month categorical variable
2 data1['month'].value_counts()
```

Out[14]:

```
may    13766
jul     6895
aug     6247
jun     5341
nov     3970
apr     2932
feb     2649
jan     1403
oct      738
sep      579
mar      477
dec      214
Name: month, dtype: int64
```

```
In [15]: 1 order={'month':{'jan':1,'feb':2,'mar':3,'apr':4,'may':5,'jun':6,'jul':7,'aug':8,'sep':9,'oct':10,'nov':11,'dec':12}}
```

```
In [16]: 1 data1=data1.replace(order)
```

```
In [17]: 1 data1
```

Out[17]:

	age	default	balance	housing	loan	day	month	duration	campaign	pdays	previous	y	job_admin.	job_blue-collar	job_entrepreneur	job_housemaid	jo
0	58	0	2143	1	0	5	5	261	1	-1	0	0	0	0	0	0	
1	44	0	29	1	0	5	5	151	1	-1	0	0	0	0	0	0	
2	33	0	2	1	1	5	5	76	1	-1	0	0	0	0	1	0	
3	47	0	1506	1	0	5	5	92	1	-1	0	0	0	1	0	0	
4	33	0	1	0	0	5	5	198	1	-1	0	0	0	0	0	0	
...
45206	51	0	825	0	0	17	11	977	3	-1	0	1	0	0	0	0	
45207	71	0	1729	0	0	17	11	456	2	-1	0	1	0	0	0	0	
45208	72	0	5715	0	0	17	11	1127	5	184	3	1	0	0	0	0	
45209	57	0	668	0	0	17	11	508	4	-1	0	0	0	1	0	0	
45210	37	0	2971	0	0	17	11	361	2	188	11	0	0	0	1	0	

45211 rows × 38 columns

```
In [18]: 1 data1.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 45211 entries, 0 to 45210
Data columns (total 38 columns):
#   Column                Non-Null Count  Dtype
---  -
0   age                    45211 non-null  int64
1   default                45211 non-null  int32
2   balance                45211 non-null  int64
3   housing                45211 non-null  int32
4   loan                   45211 non-null  int32
5   day                    45211 non-null  int64
6   month                  45211 non-null  int64
7   duration               45211 non-null  int64
8   campaign               45211 non-null  int64
9   pdays                  45211 non-null  int64
10  previous                45211 non-null  int64
11  y                       45211 non-null  int32
12  job_admin.             45211 non-null  uint8
13  job_blue-collar        45211 non-null  uint8
14  job_entrepreneur        45211 non-null  uint8
15  job_housemaid           45211 non-null  uint8
16  job_management          45211 non-null  uint8
17  job_retired             45211 non-null  uint8
18  job_self-employed       45211 non-null  uint8
19  job_services            45211 non-null  uint8
20  job_student             45211 non-null  uint8
21  job_technician          45211 non-null  uint8
22  job_unemployed          45211 non-null  uint8
23  job_unknown             45211 non-null  uint8
24  marital_divorced        45211 non-null  uint8
25  marital_married         45211 non-null  uint8
26  marital_single          45211 non-null  uint8
27  education_primary       45211 non-null  uint8
28  education_secondary     45211 non-null  uint8
29  education_tertiary      45211 non-null  uint8
30  education_unknown       45211 non-null  uint8
31  contact_cellular        45211 non-null  uint8
32  contact_telephone       45211 non-null  uint8
33  contact_unknown         45211 non-null  uint8
34  poutcome_failure        45211 non-null  uint8
35  poutcome_other          45211 non-null  uint8
36  poutcome_success        45211 non-null  uint8
37  poutcome_unknown        45211 non-null  uint8
dtypes: int32(4), int64(8), uint8(26)
memory usage: 4.6 MB
```

Model Building

Dividing our data into input and output variables

```
In [19]: 1 x=pd.concat([data1.iloc[:,0:11],data1.iloc[:,12:]],axis=1)
2 y=data1.iloc[:,11]
```

```
In [20]: 1 # Logistic regression model
2 classifier=LogisticRegression()
3 classifier.fit(x,y)
```

```
Out[20]: LogisticRegression()
```

Model Predictions

```
In [21]: 1 # Predict for x dataset
2 y_pred=classifier.predict(x)
3 y_pred
```

```
Out[21]: array([0, 0, 0, ..., 1, 0, 0])
```

```
In [22]: 1 y_pred_df=pd.DataFrame({'actual_y':y, 'y_pred_prob':y_pred})
2 y_pred_df
```

Out[22]:

	actual_y	y_pred_prob
0	0	0
1	0	0
2	0	0
3	0	0
4	0	0
...
45206	1	1
45207	1	0
45208	1	1
45209	0	0
45210	0	0

45211 rows × 2 columns

Testing Model Accuracy

```
In [23]: 1 # confusion matrix for the model Accuracy
2 confusion_matrix=confusion_matrix(y,y_pred)
3 confusion_matrix
```

Out[23]: array([[39146, 776],
[4046, 1243]], dtype=int64)


```
In [24]: 1 # The mode accuracy is calculated by (a+d)/(a+b+c+d)
2 (39156+1162)/(39156+766+4127+1162)
```

Out[24]: 0.8917741257658535

```
In [25]: 1 # As Accuracy = 0.8933, which is greater than 0.5; Thus[:,1] Threshold value>0.5=1 else[:,0] Threshold value<0.5=0
2 classifier.predict_proba(x)[: ,1]
```

Out[25]: array([0.04476732, 0.02155565, 0.01339782, ..., 0.79033147, 0.07604012,
0.12629756])

```
In [26]: 1 # ROC Curve plotting and finding AUC value
2 fpr,tpr,thresholds=roc_curve(y,classifier.predict_proba(x)[: ,1])
3 plt.plot(fpr,tpr,color='red')
4 auc=roc_auc_score(y,y_pred)
5
6 plt.plot(fpr,tpr,color='red',label='logit model(area = %0.2f)'%auc)
7 plt.plot([0,1],[0,1], 'k--')
8 plt.xlabel('False Positive Rate or [1 - True Negative Rate]')
9 plt.ylabel('True Positive Rate')
10 plt.show()
11
12 print('auc accuracy:',auc)
```



auc accuracy: 0.6077890835891568

In []:

1