

Assignment 16 (Neural_Network)(Gas_turbines)

```
In [1]: 1 import pandas as pd
2 import numpy as npd
3 import matplotlib.pyplot as plt
4 %matplotlib inline
5 import seaborn as sns
6 from sklearn.preprocessing import MinMaxScaler
7 from sklearn.preprocessing import StandardScaler
8 from sklearn.preprocessing import LabelEncoder
9 import warnings
10 warnings.filterwarnings(action='ignore')
11 from sklearn.model_selection import GridSearchCV, cross_val_score, train_test_split
12
13
14 #Load data
15 df = pd.read_csv('gas_turbines.csv')
16 df.head()
```

```
Out[1]:
```

| | AT | AP | AH | AFDP | GTEP | TIT | TAT | TEY | CDP | CO | NOX |
|---|--------|--------|--------|--------|--------|--------|--------|--------|--------|--------|--------|
| 0 | 6.8594 | 1007.9 | 96.799 | 3.5000 | 19.663 | 1059.2 | 550.00 | 114.70 | 10.605 | 3.1547 | 82.722 |
| 1 | 6.7850 | 1008.4 | 97.118 | 3.4998 | 19.728 | 1059.3 | 550.00 | 114.72 | 10.598 | 3.2363 | 82.776 |
| 2 | 6.8977 | 1008.8 | 95.939 | 3.4824 | 19.779 | 1059.4 | 549.87 | 114.71 | 10.601 | 3.2012 | 82.468 |
| 3 | 7.0569 | 1009.2 | 95.249 | 3.4805 | 19.792 | 1059.6 | 549.99 | 114.72 | 10.606 | 3.1923 | 82.670 |
| 4 | 7.3978 | 1009.7 | 95.150 | 3.4976 | 19.765 | 1059.7 | 549.98 | 114.72 | 10.612 | 3.2484 | 82.311 |

```
In [2]: 1 df.shape
```

```
Out[2]: (15039, 11)
```

```
In [3]: 1 df.columns
```

```
Out[3]: Index(['AT', 'AP', 'AH', 'AFDP', 'GTEP', 'TIT', 'TAT', 'TEY', 'CDP', 'CO',
              'NOX'],
              dtype='object')
```

```
In [4]: 1 df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 15039 entries, 0 to 15038
Data columns (total 11 columns):
#   Column  Non-Null Count  Dtype
---  ---
0    AT      15039 non-null    float64
1    AP      15039 non-null    float64
2    AH      15039 non-null    float64
3    AFDP    15039 non-null    float64
4    GTEP    15039 non-null    float64
5    TIT     15039 non-null    float64
6    TAT     15039 non-null    float64
7    TEY     15039 non-null    float64
8    CDP     15039 non-null    float64
9    CO      15039 non-null    float64
10   NOX     15039 non-null    float64
dtypes: float64(11)
memory usage: 1.3 MB
```

```
In [5]: 1 df.describe().T
```

Out[5]:

| | count | mean | std | min | 25% | 50% | 75% | max |
|------|---------|-------------|-----------|-------------|-------------|-----------|-----------|-----------|
| AT | 15039.0 | 17.764381 | 7.574323 | 0.522300 | 11.408000 | 18.1860 | 23.8625 | 34.9290 |
| AP | 15039.0 | 1013.199240 | 6.410760 | 985.850000 | 1008.900000 | 1012.8000 | 1016.9000 | 1034.2000 |
| AH | 15039.0 | 79.124174 | 13.793439 | 30.344000 | 69.750000 | 82.2660 | 90.0435 | 100.2000 |
| AFDP | 15039.0 | 4.200294 | 0.760197 | 2.087400 | 3.723900 | 4.1862 | 4.5509 | 7.6106 |
| GTEP | 15039.0 | 25.419061 | 4.173916 | 17.878000 | 23.294000 | 25.0820 | 27.1840 | 37.4020 |
| TIT | 15039.0 | 1083.798770 | 16.527806 | 1000.800000 | 1079.600000 | 1088.7000 | 1096.0000 | 1100.8000 |
| TAT | 15039.0 | 545.396183 | 7.866803 | 512.450000 | 542.170000 | 549.8900 | 550.0600 | 550.6100 |
| TEY | 15039.0 | 134.188464 | 15.829717 | 100.170000 | 127.985000 | 133.7800 | 140.8950 | 174.6100 |
| CDP | 15039.0 | 12.102353 | 1.103196 | 9.904400 | 11.622000 | 12.0250 | 12.5780 | 15.0810 |
| CO | 15039.0 | 1.972499 | 2.222206 | 0.000388 | 0.858055 | 1.3902 | 2.1604 | 44.1030 |
| NOX | 15039.0 | 68.190934 | 10.470586 | 27.765000 | 61.303500 | 66.6010 | 73.9355 | 119.8900 |

EDA & Feature Engineering

```
In [6]: 1 # Check for missing values
        2 df.isna().sum()
```


Out[6]: AT 0
AP 0
AH 0
AFDP 0
GTEP 0
TIT 0
TAT 0
TEY 0
CDP 0
CO 0
NOX 0
dtype: int64

```
In [7]: 1 df.isna().any()
```

Out[7]: AT False
AP False
AH False
AFDP False
GTEP False
TIT False
TAT False
TEY False
CDP False
CO False
NOX False
dtype: bool

```
In [8]: 1 plt.rcParams['figure.figsize']=(12,6)
        2 sns.heatmap(df.isna(), cmap =('gist_heat'), yticklabels=False)
```

Out[8]: <AxesSubplot:>



```
In [9]: 1 # check for duplicate values
        2 df[df.duplicated()].shape
```

Out[9]: (0, 11)

```
In [10]: 1 df[df.duplicated()]
```

Out[10]:

| | AT | AP | AH | AFDP | GTEP | TIT | TAT | TEY | CDP | CO | NOX |
|--|----|----|----|------|------|-----|-----|-----|-----|----|-----|
| | | | | | | | | | | | |

```
In [11]: 1 df.dtypes
```

Out[11]:

| | |
|--------|---------|
| AT | float64 |
| AP | float64 |
| AH | float64 |
| AFDP | float64 |
| GTEP | float64 |
| TIT | float64 |
| TAT | float64 |
| TEY | float64 |
| CDP | float64 |
| CO | float64 |
| NOX | float64 |
| dtype: | object |

```
In [12]: 1 df.nunique()
```

Out[12]:

| | |
|--------|-------|
| AT | 12086 |
| AP | 540 |
| AH | 12637 |
| AFDP | 11314 |
| GTEP | 8234 |
| TIT | 706 |
| TAT | 2340 |
| TEY | 4207 |
| CDP | 3611 |
| CO | 13096 |
| NOX | 11996 |
| dtype: | int64 |


Observation:

- 1.No missing values
- 2.No duplicated values
- 3.All dtypes are correct

Data Visualisation


```
In [13]: 1 # Target variable
        2 plt.title('Displot for TEY', fontsize=17, y = 1.01)
        3 sns.distplot(df['TEY'])
```

Out[13]: <AxesSubplot:title={'center':'Displot for TEY'}, xlabel='TEY', ylabel='Density'>



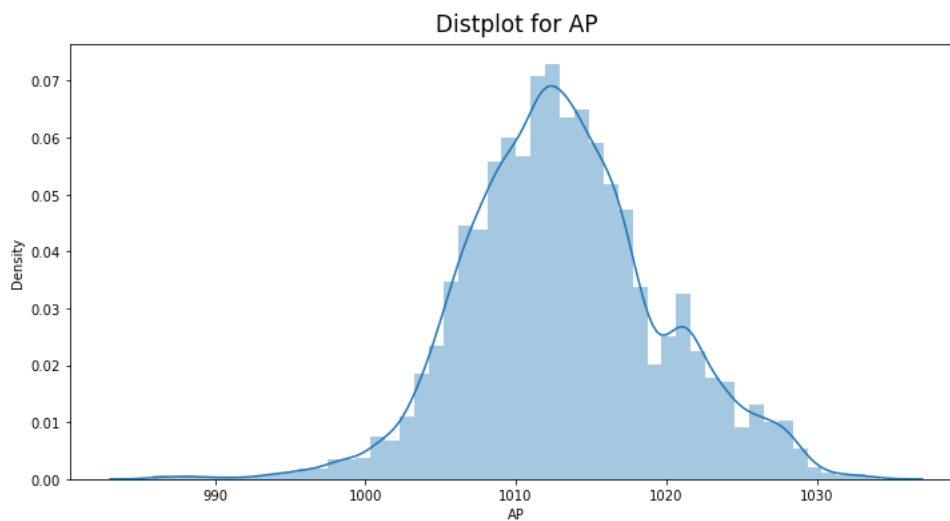
```
In [14]: 1 plt.title('Distplot for AT', fontsize=17, y = 1.01)
          2 sns.distplot(df['AT'])
```

```
Out[14]: <AxesSubplot:title={'center':'Distplot for AT'}, xlabel='AT', ylabel='Density'>
```




```
In [15]: 1 plt.title('Distplot for AP', fontsize=17, y = 1.01)
          2 sns.distplot(df['AP'])
```

```
Out[15]: <AxesSubplot:title={'center':'Distplot for AP'}, xlabel='AP', ylabel='Density'>
```



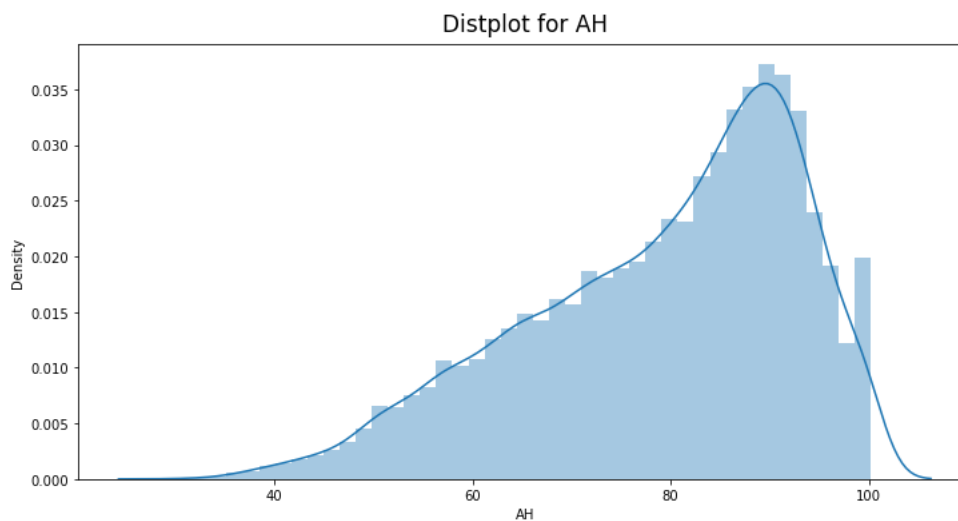
```
In [16]: 1 plt.title('Distplot for AP', fontsize=17, y = 1.01)
        2 sns.distplot(df['AP'])
```

```
Out[16]: <AxesSubplot:title={'center':'Distplot for AP'}, xlabel='AP', ylabel='Density'>
```



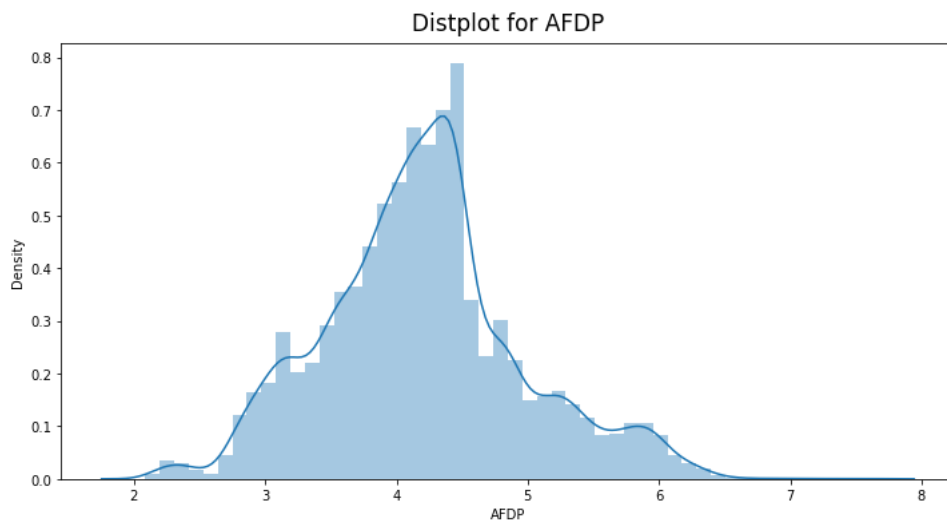
```
In [17]: 1 plt.title('Distplot for AH', fontsize=17, y = 1.01)
        2 sns.distplot(df['AH'])
```

```
Out[17]: <AxesSubplot:title={'center':'Distplot for AH'}, xlabel='AH', ylabel='Density'>
```




```
In [18]: 1 plt.title('Distplot for AFDP', fontsize=17, y = 1.01)
2         sns.distplot(df['AFDP'])
```

```
Out[18]: <AxesSubplot:title={'center':'Distplot for AFDP'}, xlabel='AFDP', ylabel='Density'>
```




```
In [19]: 1 plt.title('Distplot for GTEP', fontsize=17, y = 1.01)
2         sns.distplot(df['GTEP'])
```

```
Out[19]: <AxesSubplot:title={'center':'Distplot for GTEP'}, xlabel='GTEP', ylabel='Density'>
```



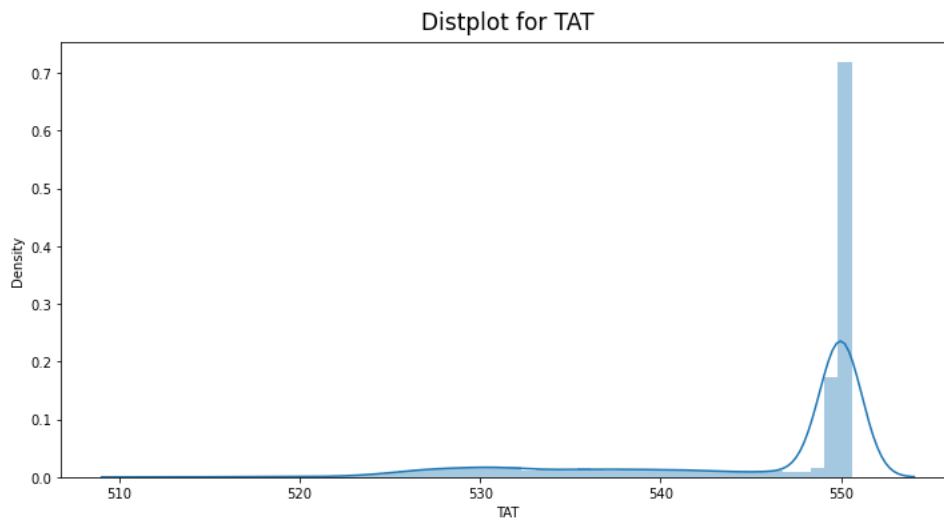
```
In [20]: 1 plt.title('Distplot for TIT', fontsize=17, y = 1.01)
        2 sns.distplot(df['TIT'])
```

```
Out[20]: <AxesSubplot:title={'center':'Distplot for TIT'}, xlabel='TIT', ylabel='Density'>
```



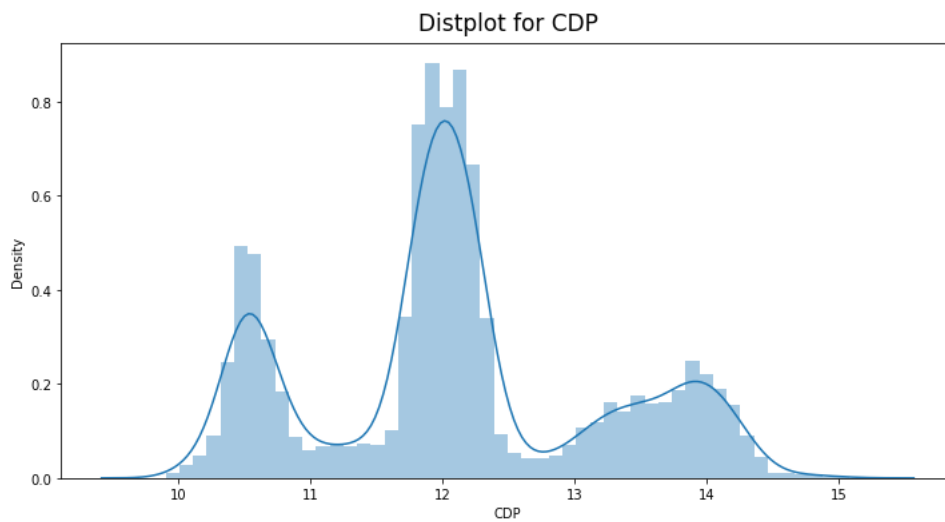
```
In [21]: 1 plt.title('Distplot for TAT', fontsize=17, y = 1.01)
        2 sns.distplot(df['TAT'])
```

```
Out[21]: <AxesSubplot:title={'center':'Distplot for TAT'}, xlabel='TAT', ylabel='Density'>
```



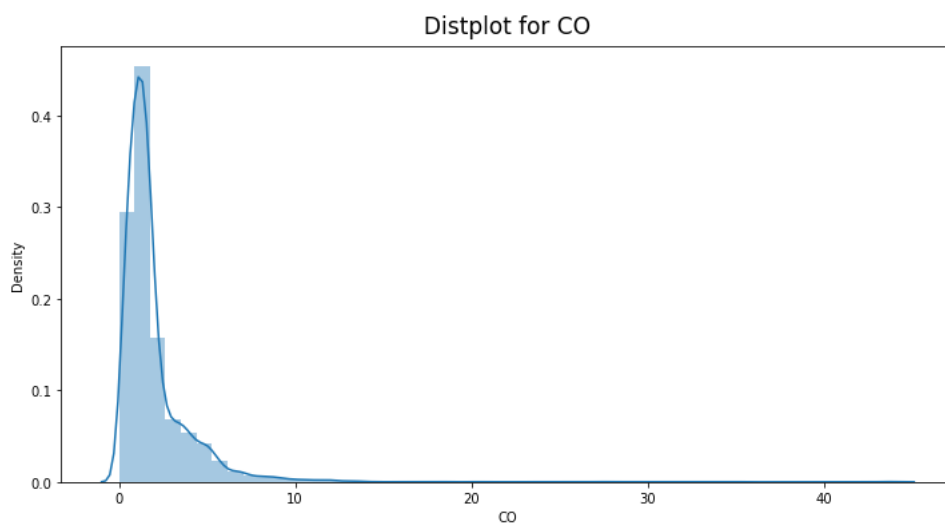
```
In [22]: 1 plt.title('Distplot for CDP', fontsize=17, y = 1.01)
        2 sns.distplot(df['CDP'])
```

```
Out[22]: <AxesSubplot:title={'center':'Distplot for CDP'}, xlabel='CDP', ylabel='Density'>
```



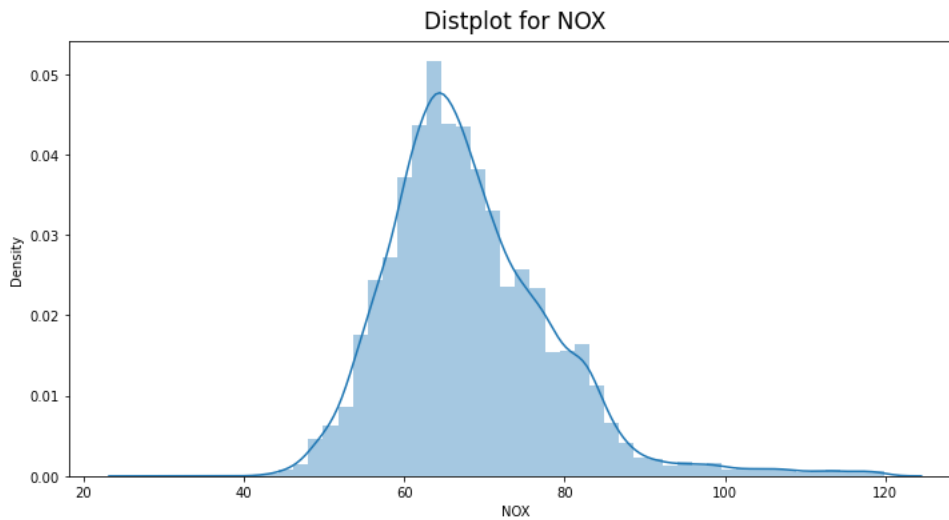
```
In [23]: 1 plt.title('Distplot for CO', fontsize=17, y = 1.01)
        2 sns.distplot(df['CO'])
```

```
Out[23]: <AxesSubplot:title={'center':'Distplot for CO'}, xlabel='CO', ylabel='Density'>
```




```
In [24]: 1 plt.title('Distplot for NOX', fontsize=17, y = 1.01)
2         sns.distplot(df['NOX'])
```

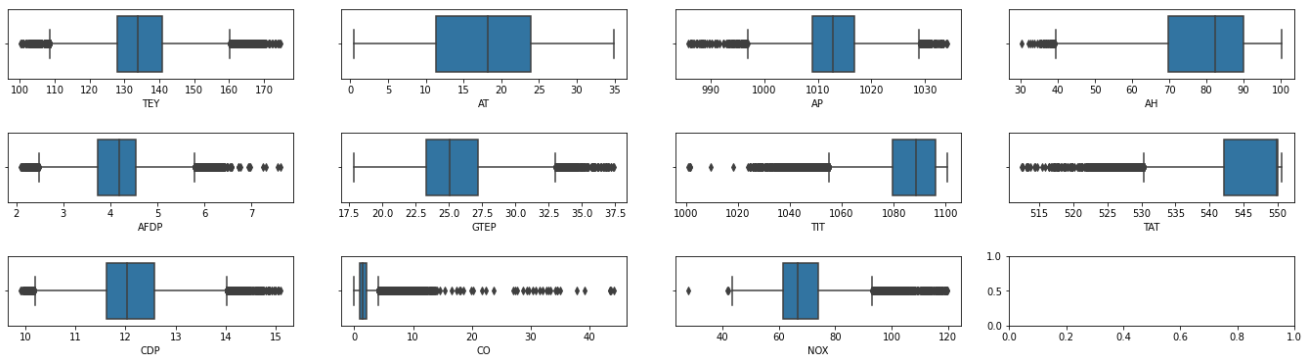
```
Out[24]: <AxesSubplot:title={'center':'Distplot for NOX'}, xlabel='NOX', ylabel='Density'>
```



```
In [25]: 1 import numpy as np
```

```
In [26]: 1 #check for outliers
2 fig, ax=plt.subplots(3,4, figsize=(19,6), sharex= False, sharey = False)
3 sns.boxplot(df.TEY, ax=ax[0,0])
4 sns.boxplot(df.AT, ax=ax[0,1])
5 sns.boxplot(df.AP, ax=ax[0,2])
6 sns.boxplot(df.AH, ax=ax[0,3])
7 sns.boxplot(df.AFD, ax=ax[1,0])
8 sns.boxplot(df.GTEP, ax=ax[1,1])
9 sns.boxplot(df.TIT, ax=ax[1,2])
10 sns.boxplot(df.TAT, ax=ax[1,3])
11 sns.boxplot(df.CDP, ax=ax[2,0])
12 sns.boxplot(df.CO, ax=ax[2,1])
13 sns.boxplot(df.NOX, ax=ax[2,2])
14 plt.suptitle("Boxplot for Continuous Variables", fontsize= 17, y = 1.06)
15 plt.tight_layout(pad=2.0)
16
```


Boxplot for Continuous Variables



1.We have a noisy data

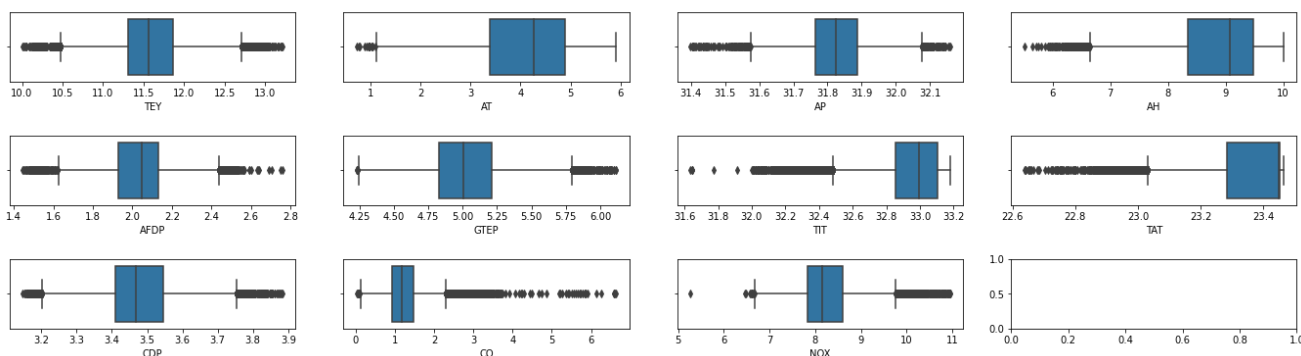
```
In [27]: 1 # check for outliers
2 fig, ax=plt.subplots(4,3, figsize=(19,6), sharex = False, sharey = False)
3 sns.boxplot(np.log(df.TEY), ax=ax[0,0])
4 sns.boxplot(np.log(df.AT), ax=ax[0,1])
5 sns.boxplot(np.log(df.AP), ax=ax[0,2])
6 sns.boxplot(np.log(df.AH), ax=ax[1,0])
7 sns.boxplot(np.log(df.AFDp), ax=ax[1,1])
8 sns.boxplot(np.log(df.GTEP), ax=ax[1,2])
9 sns.boxplot(np.log(df.TIT), ax=ax[2,0])
10 sns.boxplot(np.log(df.TAT), ax=ax[2,1])
11 sns.boxplot(np.log(df.CDP), ax=ax[2,2])
12 sns.boxplot(np.log(df.CO), ax=ax[3,0])
13 sns.boxplot(np.log(df.NOx), ax=ax[3,1])
14 plt.suptitle("Log Transformation for Continuous Variables", fontsize=17, y = 1.06)
15 plt.tight_layout(pad=2.0)
```

Log Transformation for Continuous Variables



```
In [28]: 1 fig, ax=plt.subplots(3,4, figsize=(19,6), sharex = False, sharey = False)
2 sns.boxplot(np.sqrt(df.TEY), ax=ax[0,0])
3 sns.boxplot(np.sqrt(df.AT), ax=ax[0,1])
4 sns.boxplot(np.sqrt(df.AP), ax=ax[0,2])
5 sns.boxplot(np.sqrt(df.AH), ax=ax[0,3])
6 sns.boxplot(np.sqrt(df.AFDp), ax=ax[1,0])
7 sns.boxplot(np.sqrt(df.GTEP), ax=ax[1,1])
8 sns.boxplot(np.sqrt(df.TIT), ax=ax[1,2])
9 sns.boxplot(np.sqrt(df.TAT), ax=ax[1,3])
10 sns.boxplot(np.sqrt(df.CDP), ax=ax[2,0])
11 sns.boxplot(np.sqrt(df.CO), ax=ax[2,1])
12 sns.boxplot(np.sqrt(df.NOx), ax=ax[2,2])
13 plt.suptitle("SQRT Transformation for Continuous Variables", fontsize= 17, y = 1.06)
14 plt.tight_layout(pad=2.0)
15
```

SQRT Transformation for Continuous Variables




```

In [29]: 1 fig, ax=plt.subplots(4,3, figsize=(19,6), sharex= False, sharey = False)
2 sns.boxplot(np.cbrt(df.TEY), ax=ax[0,0])
3 sns.boxplot(np.cbrt(df.AT), ax=ax[0,1])
4 sns.boxplot(np.cbrt(df.AP), ax=ax[0,2])
5 sns.boxplot(np.cbrt(df.AH), ax=ax[1,0])
6 sns.boxplot(np.cbrt(df.AFDP), ax=ax[1,1])
7 sns.boxplot(np.cbrt(df.GTEP), ax=ax[1,2])
8 sns.boxplot(np.cbrt(df.TIT), ax=ax[2,0])
9 sns.boxplot(np.cbrt(df.TAT), ax=ax[2,1])
10 sns.boxplot(np.cbrt(df.CDP), ax=ax[2,2])
11 sns.boxplot(np.cbrt(df.CO), ax=ax[3,0])
12 sns.boxplot(np.cbrt(df.NOX), ax=ax[3,1])
13 plt.suptitle("Cbrt Transformation for Continuous Variables", fontsize= 17, y = 1.06)
14 plt.tight_layout(pad=2.0)

```

Cbrt Transformation for Continuous Variables

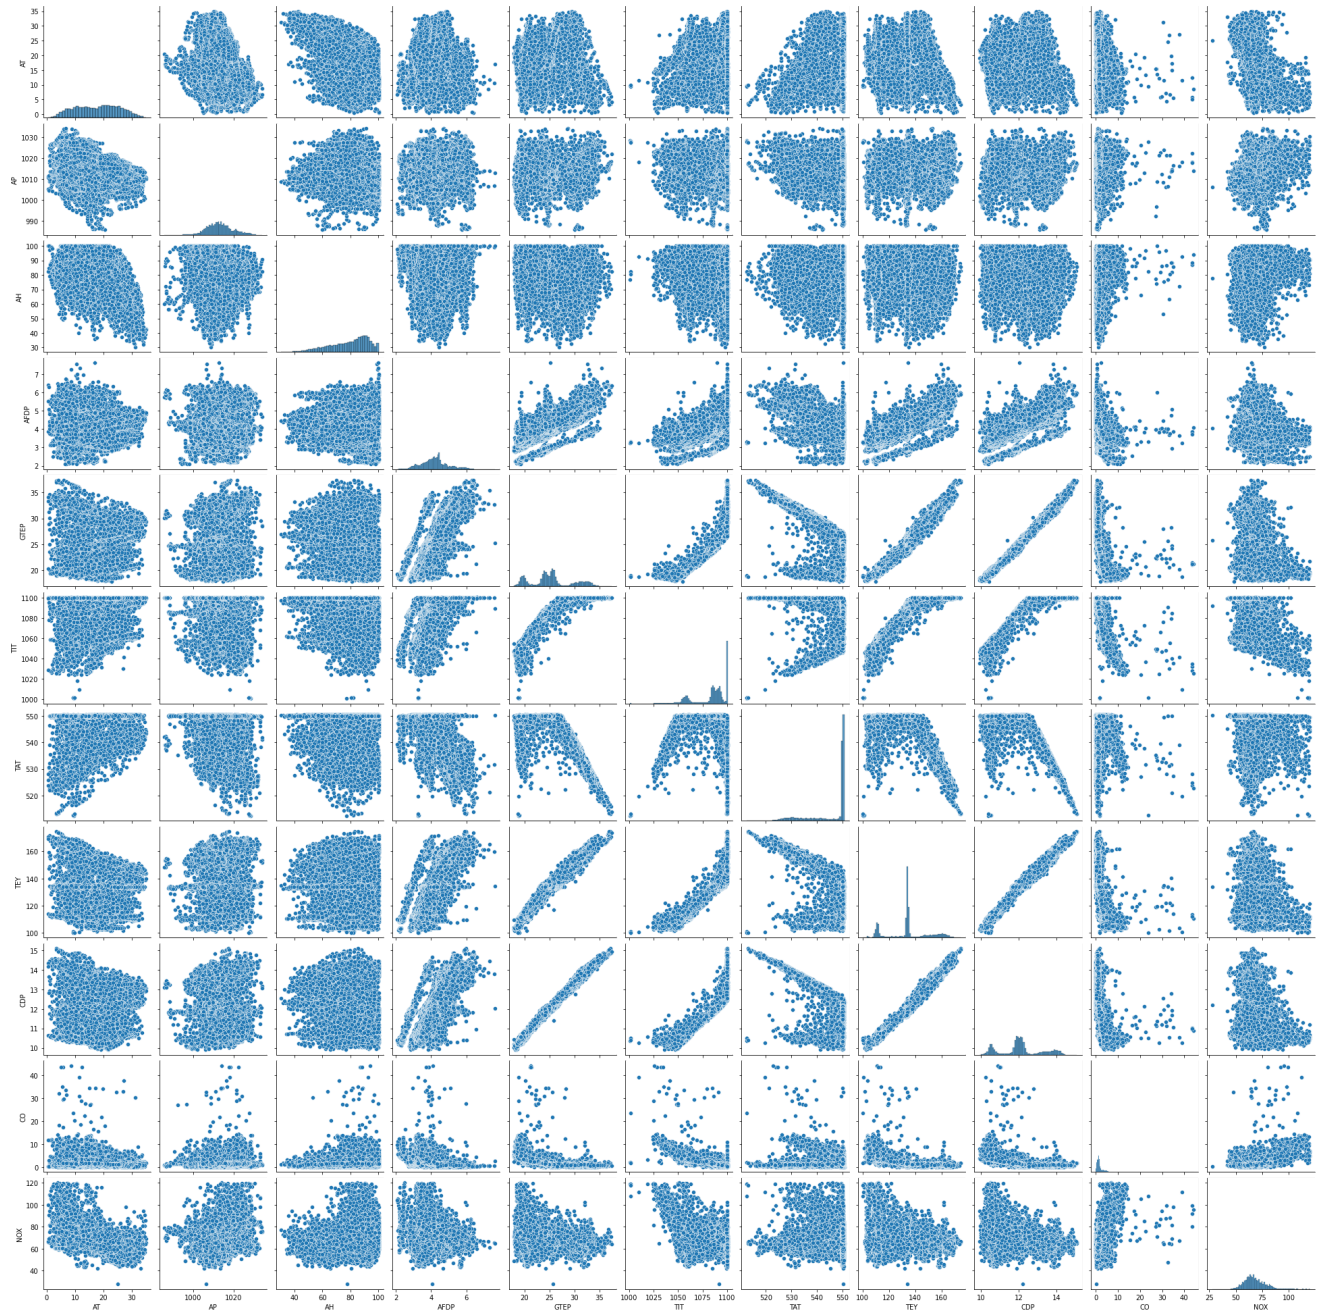


1.None of the transformations are helpful to treat the outliers.

Dependency of Target variable on diff Features.

```
In [30]: 1 sns.pairplot(df)
```

```
Out[30]: <seaborn.axisgrid.PairGrid at 0x141556a4a60>
```




```
In [31]: 1 corr = pd.DataFrame(data = df.corr().iloc[:,7], index=df.columns)
          2 corr = corr.sort_values(by='TEY', ascending=False)
          3 corr
```

Out[31]:

| | TEY |
|------|-----------|
| TEY | 1.000000 |
| CDP | 0.988473 |
| GTEP | 0.977042 |
| TIT | 0.891587 |
| AFDP | 0.717995 |
| AP | 0.146939 |
| NOX | -0.102631 |
| AH | -0.110272 |
| AT | -0.207495 |
| CO | -0.541751 |
| TAT | -0.720356 |


```
In [32]: 1 plt.title("Correlation plot between Target variables and indepedent variables",y=1.01, fontsize=18)
          2 sns.barplot(x = corr.index, y = corr.TEY)
```

Out[32]: <AxesSubplot:title={'center':'Correlation plot between Target variables and indepedent variables'}, ylabel='TEY'>



```
In [33]: 1 fig= plt.figure(figsize=(18, 10))
2         sns.heatmap(df.corr(), annot=True);
3         plt.xticks(rotation=45)
4         plt.title("Correlation Map of Variables", fontsize=19)
```

Out[33]: Text(0.5, 1.0, 'Correlation Map of Variables')



```
In [34]: 1 !pip install ppscore
```

Requirement already satisfied: ppscore in c:\users\admin\anaconda3\lib\site-packages (1.3.0)
Requirement already satisfied: pandas<2.0.0,>=1.0.0 in c:\users\admin\anaconda3\lib\site-packages (from ppscore) (1.3.4)
Requirement already satisfied: scikit-learn<2.0.0,>=0.20.2 in c:\users\admin\anaconda3\lib\site-packages (from ppscore) (1.1.3)
Requirement already satisfied: pytz>=2017.3 in c:\users\admin\anaconda3\lib\site-packages (from pandas<2.0.0,>=1.0.0->ppscore) (2021.3)
Requirement already satisfied: numpy>=1.17.3 in c:\users\admin\anaconda3\lib\site-packages (from pandas<2.0.0,>=1.0.0->ppscore) (1.23.5)
Requirement already satisfied: python-dateutil>=2.7.3 in c:\users\admin\anaconda3\lib\site-packages (from pandas<2.0.0,>=1.0.0->ppscore) (2.8.2)
Requirement already satisfied: six>=1.5 in c:\users\admin\anaconda3\lib\site-packages (from python-dateutil>=2.7.3->pandas<2.0.0,>=1.0.0->ppscore) (1.16.0)
Requirement already satisfied: scipy>=1.3.2 in c:\users\admin\anaconda3\lib\site-packages (from scikit-learn<2.0.0,>=0.20.2->ppscore) (1.9.3)
Requirement already satisfied: joblib>=1.0.0 in c:\users\admin\anaconda3\lib\site-packages (from scikit-learn<2.0.0,>=0.20.2->ppscore) (1.1.0)
Requirement already satisfied: threadpoolctl>=2.0.0 in c:\users\admin\anaconda3\lib\site-packages (from scikit-learn<2.0.0,>=0.20.2->ppscore) (2.2.0)


```
In [35]: 1 import ppscore as PPS
2 score = PPS.matrix(df)
3 score_s = score[score['y']=='TEY']
4 score_s.sort_values(by="ppscore", ascending=False)
```

```
Out[35]:
```

| | x | y | ppscore | case | is_valid_score | metric | baseline_score | model_score | model |
|-----|------|-----|----------|----------------|----------------|---------------------|----------------|-------------|-------------------------|
| 84 | TEY | TEY | 1.000000 | predict_itself | True | None | 0.000000 | 1.000000 | None |
| 95 | CDP | TEY | 0.872285 | regression | True | mean absolute error | 11.172076 | 1.426840 | DecisionTreeRegressor() |
| 51 | GTEP | TEY | 0.832336 | regression | True | mean absolute error | 11.172076 | 1.873154 | DecisionTreeRegressor() |
| 62 | TIT | TEY | 0.766040 | regression | True | mean absolute error | 11.172076 | 2.613821 | DecisionTreeRegressor() |
| 73 | TAT | TEY | 0.226050 | regression | True | mean absolute error | 11.172076 | 8.646631 | DecisionTreeRegressor() |
| 40 | AFDP | TEY | 0.152509 | regression | True | mean absolute error | 11.172076 | 9.468234 | DecisionTreeRegressor() |
| 106 | CO | TEY | 0.055869 | regression | True | mean absolute error | 11.172076 | 10.547906 | DecisionTreeRegressor() |
| 7 | AT | TEY | 0.000000 | regression | True | mean absolute error | 11.172076 | 16.007470 | DecisionTreeRegressor() |
| 18 | AP | TEY | 0.000000 | regression | True | mean absolute error | 11.172076 | 12.475617 | DecisionTreeRegressor() |
| 29 | AH | TEY | 0.000000 | regression | True | mean absolute error | 11.172076 | 16.950976 | DecisionTreeRegressor() |
| 117 | NOX | TEY | 0.000000 | regression | True | mean absolute error | 11.172076 | 14.537337 | DecisionTreeRegressor() |

```
In [36]: 1 plt.rcParams['figure.figsize']=(19,6)
2 sns.barplot(x='x', y='ppscore', data=score_s.sort_values(by='ppscore', ascending=False))
3 plt.title("PPScore of each feature with Target various", fontsize=17, y=1.01)
```

```
Out[36]: Text(0.5, 1.01, 'PPScore of each feature with Target various')
```



Observation:

1. From correlation matrix as well as ppscore we can clearly see that TEY is highly dependent on 'CDP', 'GTEP', 'TIT'.
2. We can drop 'AT', 'AP', 'AH', 'NOX', as they have very less impact on dependent variables.

Check for outliers.

```
In [37]: 1 # Check for Outliers
2 from sklearn.ensemble import IsolationForest
3 data1=df.copy()
4
5 #training the model
6 clf = IsolationForest(random_state=10, contamination=.001)
7 clf.fit(data1)
8 data1['anomaly'] = clf.predict(data1.iloc[:,0:11])
9 outliers = data1[data1['anomaly']==-1]
```

```
In [38]: 1 outliers
```

Out[38]:

| | AT | AP | AH | AFDP | GTEP | TIT | TAT | TEY | CDP | CO | NOX | anamoly |
|-------|----------|---------|--------|--------|--------|--------|--------|--------|--------|---------|---------|---------|
| 261 | 5.66020 | 1018.30 | 86.968 | 3.8404 | 21.079 | 1028.5 | 523.86 | 112.02 | 10.963 | 43.4280 | 99.237 | -1 |
| 553 | 3.55320 | 1027.30 | 90.871 | 4.2162 | 21.464 | 1041.2 | 531.68 | 117.76 | 10.984 | 8.8254 | 106.840 | -1 |
| 763 | 1.81300 | 1007.20 | 74.980 | 3.6967 | 19.958 | 1026.4 | 528.18 | 111.72 | 10.553 | 12.0900 | 114.940 | -1 |
| 764 | 1.49880 | 1006.30 | 76.734 | 3.7063 | 20.041 | 1027.6 | 528.79 | 112.28 | 10.585 | 11.6520 | 112.830 | -1 |
| 765 | 0.97877 | 1005.70 | 78.978 | 3.7379 | 20.084 | 1027.9 | 528.52 | 112.71 | 10.628 | 11.6910 | 108.880 | -1 |
| 993 | 4.36570 | 1021.60 | 85.528 | 3.9574 | 20.263 | 1025.6 | 525.72 | 111.35 | 10.652 | 12.7860 | 112.270 | -1 |
| 6896 | 17.13200 | 1010.80 | 80.503 | 2.2148 | 18.484 | 1034.1 | 539.98 | 102.07 | 10.182 | 11.5150 | 110.760 | -1 |
| 7019 | 7.02760 | 997.23 | 97.761 | 2.0992 | 19.227 | 1037.2 | 538.53 | 109.63 | 10.338 | 11.0440 | 105.060 | -1 |
| 7470 | 7.04730 | 1019.60 | 96.885 | 2.4558 | 19.501 | 1032.0 | 532.32 | 109.21 | 10.567 | 11.3740 | 112.230 | -1 |
| 9920 | 15.17900 | 1017.60 | 71.630 | 2.7816 | 18.435 | 1027.8 | 533.45 | 103.64 | 10.143 | 12.1440 | 113.800 | -1 |
| 13820 | 14.18300 | 1023.10 | 78.110 | 3.1557 | 18.869 | 1025.0 | 530.16 | 103.80 | 10.340 | 13.3130 | 116.340 | -1 |
| 13921 | 11.58500 | 1018.20 | 92.751 | 3.2518 | 18.784 | 1009.5 | 519.71 | 100.83 | 10.253 | 39.0500 | 111.780 | -1 |
| 14100 | 9.40970 | 1027.90 | 82.224 | 3.3003 | 18.987 | 1001.4 | 512.60 | 100.32 | 10.495 | 23.6290 | 107.890 | -1 |
| 14278 | 9.90780 | 1026.10 | 65.923 | 3.3126 | 19.366 | 1024.5 | 527.21 | 108.08 | 10.506 | 20.2710 | 105.660 | -1 |
| 14317 | 3.93850 | 1021.30 | 90.536 | 3.4765 | 20.031 | 1026.6 | 526.30 | 111.70 | 10.683 | 14.0350 | 114.700 | -1 |
| 14320 | 3.49070 | 1020.80 | 91.519 | 3.5309 | 20.098 | 1025.8 | 525.35 | 111.91 | 10.761 | 11.9210 | 113.900 | -1 |

1. These are the outlires in our data.

Data Preprocessing

```
In [39]: 1 df.shape
```

Out[39]: (15039, 11)

```
In [40]: 1 # drop the outliers
2 df = df.drop(outliers.index)
3 df.shape
```

Out[40]: (15023, 11)

```
In [41]: 1 #reset index after dropping outlires
2 df = df.reset_index()
3 df = df.drop('index', axis = 1)
4 df
```

Out[41]:

| | AT | AP | AH | AFDP | GTEP | TIT | TAT | TEY | CDP | CO | NOX |
|-------|--------|--------|--------|--------|--------|--------|--------|--------|--------|--------|--------|
| 0 | 6.8594 | 1007.9 | 96.799 | 3.5000 | 19.663 | 1059.2 | 550.00 | 114.70 | 10.605 | 3.1547 | 82.722 |
| 1 | 6.7850 | 1008.4 | 97.118 | 3.4998 | 19.728 | 1059.3 | 550.00 | 114.72 | 10.598 | 3.2363 | 82.776 |
| 2 | 6.8977 | 1008.8 | 95.939 | 3.4824 | 19.779 | 1059.4 | 549.87 | 114.71 | 10.601 | 3.2012 | 82.468 |
| 3 | 7.0569 | 1009.2 | 95.249 | 3.4805 | 19.792 | 1059.6 | 549.99 | 114.72 | 10.606 | 3.1923 | 82.670 |
| 4 | 7.3978 | 1009.7 | 95.150 | 3.4976 | 19.765 | 1059.7 | 549.98 | 114.72 | 10.612 | 3.2484 | 82.311 |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 15018 | 9.0301 | 1005.6 | 98.460 | 3.5421 | 19.164 | 1049.7 | 546.21 | 111.61 | 10.400 | 4.5186 | 79.559 |
| 15019 | 7.8879 | 1005.9 | 99.093 | 3.5059 | 19.414 | 1046.3 | 543.22 | 111.78 | 10.433 | 4.8470 | 79.917 |
| 15020 | 7.2647 | 1006.3 | 99.496 | 3.4770 | 19.530 | 1037.7 | 537.32 | 110.19 | 10.483 | 7.9632 | 90.912 |
| 15021 | 7.0060 | 1006.8 | 99.008 | 3.4486 | 19.377 | 1043.2 | 541.24 | 110.74 | 10.533 | 6.2494 | 93.227 |
| 15022 | 6.9279 | 1007.2 | 97.533 | 3.4275 | 19.306 | 1049.9 | 545.85 | 111.58 | 10.583 | 4.9816 | 92.498 |

15023 rows × 11 columns

```
In [42]: 1 df = df.drop(['AT', 'AP', 'AH', 'NOX'], axis=1)
```

```
In [43]: 1 df.shape
```

Out[43]: (15023, 7)

Converting independent feature into normalised and standardized data

```
In [44]: 1 # Standardize & Normalize the data
2 norm = MinMaxScaler()
3 std = StandardScaler()
4
5 df_norm = pd.DataFrame(norm.fit_transform(df), columns=df.columns) # data between -3 to +3
6 df_std = pd.DataFrame(std.fit_transform(df), columns=df.columns) # data between -1 to +1
```

Take a smaller sample to build a model.

```
In [45]: 1 # We will take a small model as this is a large data and will take huge ammount of time to
2 # build model to randomly shuffle and select a % of data
3 temp = df_std.sample(frac=1) # shuffle all the data
4 temp_s = df_std.sample(frac=0.1) # shuffle and select only 10% of the data randomly to train
```

```
In [46]: 1 temp_s
```

```
Out[46]:
```

| | AFDP | GTEP | TIT | TAT | TEY | CDP | CO |
|-------|-----------|-----------|-----------|-----------|-----------|-----------|-----------|
| 10473 | 0.351488 | 1.290082 | 0.976474 | -1.174103 | 1.191558 | 1.313317 | -0.509066 |
| 3911 | -0.811922 | -0.997964 | -1.196849 | 0.554989 | -1.320629 | -1.038528 | 1.564372 |
| 12188 | -0.463255 | 0.081457 | 0.489455 | 0.633989 | -0.014266 | 0.184105 | -0.774782 |
| 6408 | -0.638576 | -1.318458 | -1.568201 | 0.594489 | -1.491986 | -1.435792 | 1.329205 |
| 4776 | 0.448757 | 0.260282 | 0.629473 | 0.602134 | 0.025569 | 0.116987 | -0.612006 |
| ... | ... | ... | ... | ... | ... | ... | ... |
| 392 | 0.471659 | -0.247427 | 0.081576 | 0.575376 | 0.001541 | -0.169624 | -0.014359 |
| 1909 | 0.040333 | -1.037516 | -1.135972 | 0.580473 | -1.019015 | -1.050318 | 0.158078 |
| 11709 | 0.365572 | 1.104546 | 0.982562 | -0.686084 | 0.862754 | 1.008565 | -0.857748 |
| 4850 | -0.366907 | -0.560251 | -0.423706 | 0.546069 | -0.712974 | -0.522446 | -0.262947 |
| 9785 | 0.415062 | 1.608898 | 0.976474 | -1.882560 | 1.651251 | 1.626231 | -0.492619 |

1502 rows × 7 columns

Splitting data into target variable and independent variables.

```
In [47]: 1 x = temp_s.drop('TEY', axis=1)
2 y = temp_s['TEY']
3 x
```

```
Out[47]:
```

| | AFDP | GTEP | TIT | TAT | CDP | CO |
|-------|-----------|-----------|-----------|-----------|-----------|-----------|
| 10473 | 0.351488 | 1.290082 | 0.976474 | -1.174103 | 1.313317 | -0.509066 |
| 3911 | -0.811922 | -0.997964 | -1.196849 | 0.554989 | -1.038528 | 1.564372 |
| 12188 | -0.463255 | 0.081457 | 0.489455 | 0.633989 | 0.184105 | -0.774782 |
| 6408 | -0.638576 | -1.318458 | -1.568201 | 0.594489 | -1.435792 | 1.329205 |
| 4776 | 0.448757 | 0.260282 | 0.629473 | 0.602134 | 0.116987 | -0.612006 |
| ... | ... | ... | ... | ... | ... | ... |
| 392 | 0.471659 | -0.247427 | 0.081576 | 0.575376 | -0.169624 | -0.014359 |
| 1909 | 0.040333 | -1.037516 | -1.135972 | 0.580473 | -1.050318 | 0.158078 |
| 11709 | 0.365572 | 1.104546 | 0.982562 | -0.686084 | 1.008565 | -0.857748 |
| 4850 | -0.366907 | -0.560251 | -0.423706 | 0.546069 | -0.522446 | -0.262947 |
| 9785 | 0.415062 | 1.608898 | 0.976474 | -1.882560 | 1.626231 | -0.492619 |

1502 rows × 6 columns

Creating train and test data for model validation

```
In [48]: 1 x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.25)
```

```
In [49]: 1 x_train.shape, y_test.shape, y_train.shape, y_test.shape
```

```
Out[49]: ((1126, 6), (376,)), ((1126,)), (376,))
```

Build a model.

```
In [50]: 1 # Importing the neccessary packages
2 import tensorflow as tf
3 import keras
4 from sklearn.model_selection import GridSearchCV
5 from keras.models import Sequential
6 from keras.layers import Dense
7 from keras.wrappers.scikit_learn import KerasRegressor
8 from tensorflow.keras.optimizers import Adam
9 from keras.layers import Dropout
10 tf.config.experimental.list_logical_devices('GPU') # to use GPU for faster processing of model
```

```
Out[50]: []
```

```
In [51]: 1 # create model with 2 hidden layers
2 def create_model_two_hidden_layers():
3     model = Sequential()
4     model.add(Dense(5, input_dim=6, kernel_initializer='uniform', activation='relu'))
5     model.add(Dense(6, kernel_initializer='uniform', activation='relu'))
6     model.add(Dense(10, kernel_initializer='uniform', activation='relu'))
7     model.add(Dense(1))
8
9     adam=Adam(lr=0.001)
10    model.compile(loss='mse', optimizer=adam, metrics=['mse', 'mae', 'mape'])
11    return model
```

```
In [52]: 1 model1 = create_model_two_hidden_layers()
2 print("Here is the summary of the model:")
3 model1.summary()
```

Here is the summary of the model:

Model: "sequential"

| Layer (type) | Output Shape | Param # |
|-------------------------|--------------|---------|
| ===== | ===== | ===== |
| dense (Dense) | (None, 5) | 35 |
| dense_1 (Dense) | (None, 6) | 36 |
| dense_2 (Dense) | (None, 10) | 70 |
| dense_3 (Dense) | (None, 1) | 11 |
| ===== | ===== | ===== |
| Total params: 152 | | |
| Trainable params: 152 | | |
| Non-trainable params: 0 | | |

```
In [53]: 1 # create model with 3 hidden layers
2 def create_model_three_hidden_layers():
3     model = Sequential()
4     model.add(Dense(32, input_dim=6, kernel_initializer='uniform', activation='relu'))
5     model.add(Dense(32, kernel_initializer='uniform', activation='relu'))
6     model.add(Dense(64, kernel_initializer='uniform', activation='relu'))
7     model.add(Dense(128, kernel_initializer='uniform', activation='relu'))
8     model.add(Dense(1))
9
10    adam=Adam(lr=0.01)
11    model.compile(loss='mse', optimizer=adam, metrics=['mse', 'mae', 'mape'])
12    return model
```

```
In [54]: 1 model2 = create_model_three_hidden_layers()
2 print("Here is the summary of the model:")
3 model2.summary()
```

Here is the summary of the model:
Model: "sequential_1"

| Layer (type) | Output Shape | Param # |
|--------------------------|--------------|---------|
| dense_4 (Dense) | (None, 32) | 224 |
| dense_5 (Dense) | (None, 32) | 1056 |
| dense_6 (Dense) | (None, 64) | 2112 |
| dense_7 (Dense) | (None, 128) | 8320 |
| dense_8 (Dense) | (None, 1) | 129 |
| Total params: 11,841 | | |
| Trainable params: 11,841 | | |
| Non-trainable params: 0 | | |

```
In [55]: 1 %%time
2 epochs=500
3 batch_size=50
4
5 print("Here is the summary of this model:")
6 model2.summary()
7
8 with tf.device('/GPU:0'):
9     model2.fit(x_train,y_train, verbose = 0,batch_size = batch_size,epochs = epochs, shuffle=True)
```

Here is the summary of this model:
Model: "sequential_1"

| Layer (type) | Output Shape | Param # |
|--------------------------|--------------|---------|
| dense_4 (Dense) | (None, 32) | 224 |
| dense_5 (Dense) | (None, 32) | 1056 |
| dense_6 (Dense) | (None, 64) | 2112 |
| dense_7 (Dense) | (None, 128) | 8320 |
| dense_8 (Dense) | (None, 1) | 129 |
| Total params: 11,841 | | |
| Trainable params: 11,841 | | |
| Non-trainable params: 0 | | |
| Wall time: 22.2 s | | |

```
In [56]: 1 print("Predicted Values:")
2 model2.predict(x_test[:10])
```

Predicted Values:
1/1 [=====] - 0s 129ms/step

```
Out[56]: array([[ -1.833806 ],
[ -1.4910991 ],
[  0.00822169],
[ -0.00512612],
[ -1.605835  ],
[  0.9588574 ],
[ -1.4259658 ],
[  0.01884493],
[ -0.01247218],
[ -0.04152176]], dtype=float32)
```

```
In [57]: 1 print('Actual Values')
        2 y_test[:10]
```

Actual Values

```
Out[57]: 2317    -1.931445
        6485    -1.474914
        9091    -0.004782
        5443    -0.028810
        10027   -1.484398
        11794    0.800787
        10773   -1.402830
        8915     0.098286
        6588    -0.035133
        1384    -0.095203
        Name: TEY, dtype: float64
```

```
In [58]: 1 loss, mae, mse, mape = model2.evaluate(x_train, y_train)
        2 print('\n', "Results for model 2:", '\n', "Training Loss:", loss, '\n', "Training Mean Absolute Error:" , mae, '\n', "Traini
```

36/36 [=====] - 0s 2ms/step - loss: 0.0049 - mse: 0.0049 - mae: 0.0501 - mape: 76.3284

Results for model 2:
 Training Loss: 0.0049224658869206905
 Training Mean Absolute Error: 0.0049224658869206905
 Training Mean Squared Error: 0.050090059638023376

```
In [59]: 1 loss, mae, mse, mape = model2.evaluate(x_train, y_train)
        2 print('\n', "Results for model 2:", '\n', "Test Loss:", loss, '\n', "Test Mean Absolute Error:" , mae, '\n', "Test Mean Squa
```

36/36 [=====] - 0s 2ms/step - loss: 0.0049 - mse: 0.0049 - mae: 0.0501 - mape: 76.3284

Results for model 2:
 Test Loss: 0.0049224658869206905
 Test Mean Absolute Error: 0.0049224658869206905
 Test Mean Squared Error: 0.050090059638023376

Observations

1. We got pretty good results for their model.
2. Train and test errors are also quiet similar, which means our model is not overfitted or underfitted.
3. Still we will try to get best results by doing hyperparameter tuning.

Hyperparameter Tuning to get best options for:

1. batchsize
2. epochs
3. neurons
4. learning rate
5. dropout
6. kernel initializer
7. activation function

```
In [60]: 1 from numpy import array
        2 from sklearn.model_selection import KFold
```

```
In [61]: 1 # Create the model
2 #get best value for batch size and epochs by hyperparameter tuning
3 model = KerasRegressor(build_fn = create_model_three_hidden_layers,verbose = 0)
4 # Define the grid search parameters
5 batch_size = [30,50,70]
6 epochs = [300,500,800]
7 # Make a dictionary of the grid search parameters
8 param_grid = dict(batch_size = batch_size,epochs = epochs)
9 # Build and fit the GridSearchCV
10 grid = GridSearchCV(estimator = model,param_grid = param_grid,cv = KFold(),verbose = 10)
11 grid_result = grid.fit(x_train,y_train)
```

Fitting 5 folds for each of 9 candidates, totalling 45 fits

```
[CV 1/5; 1/9] START batch_size=30, epochs=300.....
[CV 1/5; 1/9] END ...batch_size=30, epochs=300;; score=-0.007 total time= 16.3s
[CV 2/5; 1/9] START batch_size=30, epochs=300.....
[CV 2/5; 1/9] END ...batch_size=30, epochs=300;; score=-0.012 total time= 17.3s
[CV 3/5; 1/9] START batch_size=30, epochs=300.....
[CV 3/5; 1/9] END ...batch_size=30, epochs=300;; score=-0.008 total time= 17.1s
[CV 4/5; 1/9] START batch_size=30, epochs=300.....
[CV 4/5; 1/9] END ...batch_size=30, epochs=300;; score=-0.007 total time= 17.2s
[CV 5/5; 1/9] START batch_size=30, epochs=300.....
[CV 5/5; 1/9] END ...batch_size=30, epochs=300;; score=-0.006 total time= 17.9s
[CV 1/5; 2/9] START batch_size=30, epochs=500.....
[CV 1/5; 2/9] END ...batch_size=30, epochs=500;; score=-0.015 total time= 26.5s
[CV 2/5; 2/9] START batch_size=30, epochs=500.....
[CV 2/5; 2/9] END ...batch_size=30, epochs=500;; score=-0.010 total time= 28.6s
[CV 3/5; 2/9] START batch_size=30, epochs=500.....
[CV 3/5; 2/9] END ...batch_size=30, epochs=500;; score=-0.010 total time= 28.7s
[CV 4/5; 2/9] START batch_size=30, epochs=500.....
[CV 4/5; 2/9] END ...batch_size=30, epochs=500;; score=-0.008 total time= 28.7s
```

```
In [62]: 1 # Summarize the results
2 print('Best : {}, using {}'.format(grid_result.best_score_,grid_result.best_params_))
3 means = grid_result.cv_results_['mean_test_score']
4 stds = grid_result.cv_results_['std_test_score']
5 params = grid_result.cv_results_['params']
6 for means, stdev, param in zip(means, stds, params):
7     print('{} with: {}'.format(means, stdev, param))
```

```
Best : -0.007391261588782072, using {'batch_size': 70, 'epochs': 300}
-0.007894857414066791,0.001991391197786599 with: {'batch_size': 30, 'epochs': 300}
-0.009753851965069772,0.0030524625775304425 with: {'batch_size': 30, 'epochs': 500}
-0.007557791378349066,0.002037569500462883 with: {'batch_size': 30, 'epochs': 800}
-0.010061297938227654,0.0044086401658283494 with: {'batch_size': 50, 'epochs': 300}
-0.009114649146795273,0.0015327226797154107 with: {'batch_size': 50, 'epochs': 500}
-0.009040820598602294,0.0027362051480239943 with: {'batch_size': 50, 'epochs': 800}
-0.007391261588782072,0.001285354873909179 with: {'batch_size': 70, 'epochs': 300}
-0.009143414068967104,0.0022328322340198397 with: {'batch_size': 70, 'epochs': 500}
-0.008697202522307634,0.0010586759703553783 with: {'batch_size': 70, 'epochs': 800}
```

```

In [63]: 1 #get best value for learning rate and dropout by hyperparameter tuning
2
3 # Defining the model
4 %time
5 def create_model_three_hidden_layers(learning_rate,dropout_rate):
6     model = Sequential()
7     model.add(Dense(32,input_dim = 6,kernel_initializer = 'uniform',activation = 'relu'))
8     model.add(Dropout(dropout_rate))
9     model.add(Dense(32,kernel_initializer = 'uniform',activation = 'relu'))
10    model.add(Dropout(dropout_rate))
11    model.add(Dense(64,kernel_initializer = 'uniform',activation = 'relu'))
12    model.add(Dropout(dropout_rate))
13    model.add(Dense(128,kernel_initializer = 'uniform',activation = 'relu'))
14    model.add(Dropout(dropout_rate))
15    model.add(Dense(1))
16
17    adam = Adam(lr = learning_rate)
18    model.compile(loss = 'mse', optimizer = adam,metrics = ['mse', 'mae', 'mape'])
19    return model
20
21 # Create the model
22
23 model = KerasRegressor(build_fn = create_model_three_hidden_layers,verbose = 0,
24                        batch_size = 70,epochs = 300)
25
26 # Define the grid search parameters
27
28 learning_rate = [0.001,0.01,0.1]
29 dropout_rate = [0.0,0.1,0.2]
30
31 # Make a dictionary of the grid search parameters
32
33 param_grids = dict(learning_rate = learning_rate,dropout_rate = dropout_rate)
34
35 # Build and fit the GridSearchCV
36
37 grid = GridSearchCV(estimator = model,param_grid = param_grids,cv = KFold(),verbose = 0)
38 grid_result = grid.fit(x_train,y_train)

```

Wall time: 0 ns

```

In [64]: 1 # Summarize the results
2 print('Best : {}, using {}'.format(grid_result.best_score_,grid_result.best_params_))
3 means = grid_result.cv_results_['mean_test_score']
4 stds = grid_result.cv_results_['std_test_score']
5 params = grid_result.cv_results_['params']
6 for mean, stdev, param in zip(means, stds, params):
7     print('{} with: {}'.format(mean, stdev, param))

```

Best : -0.007342452276498079, using {'dropout_rate': 0.0, 'learning_rate': 0.001}

-0.007342452276498079,0.0019959824416733475 with: {'dropout_rate': 0.0, 'learning_rate': 0.001}

-0.008561298809945583,0.001991463301998141 with: {'dropout_rate': 0.0, 'learning_rate': 0.01}

-0.42658794578164816,0.48784841583609684 with: {'dropout_rate': 0.0, 'learning_rate': 0.1}

-0.010010520182549953,0.0011587275738232285 with: {'dropout_rate': 0.1, 'learning_rate': 0.001}

-0.010904456581920385,0.0023298427190362866 with: {'dropout_rate': 0.1, 'learning_rate': 0.01}

-0.5191736936569213,0.36482949841031104 with: {'dropout_rate': 0.1, 'learning_rate': 0.1}

-0.011447767540812493,0.0033028640505105367 with: {'dropout_rate': 0.2, 'learning_rate': 0.001}

-0.017722824588418007,0.004725459057919966 with: {'dropout_rate': 0.2, 'learning_rate': 0.01}

-0.7296398758888245,0.28865086924640976 with: {'dropout_rate': 0.2, 'learning_rate': 0.1}

```

In [89]: 1 # Defining the model
2 #get best value for kernel initializer and activation func by hyperparameter tuning
3
4 %%time
5 def create_model_three_hidden_layers(activation_function,init):
6     model = Sequential()
7     model.add(Dense(32,input_dim = 6,kernel_initializer = init,activation = activation_function))
8
9     model.add(Dense(32,kernel_initializer = init,activation = activation_function))
10
11     model.add(Dense(64,kernel_initializer = init,activation = activation_function))
12
13     model.add(Dense(128,kernel_initializer = init,activation = activation_function))
14
15     model.add(Dense(1))
16
17     adam = Adam(lr = 0.001)
18     model.compile(loss = 'mse',optimizer = adam,metrics = ['mse', 'mae', 'mape'])
19     return model
20
21 # Create the model
22
23 model = KerasRegressor(build_fn = create_model_three_hidden_layers,verbose = 0,batch_size = 70,epochs = 300)
24
25 # Define the grid search parameters
26 activation_function = ['softmax','relu','tanh','linear']
27 init = ['uniform','normal','zero']
28
29 # Make a dictionary of the grid search parameters
30 param_grids = dict(activation_function = activation_function,init = init)
31
32 # Build and fit the GridSearchCV
33
34 grid = GridSearchCV(estimator = model,param_grid = param_grids,cv = KFold(),verbose = 0)
35 grid_result = grid.fit(x_train,y_train)

```

UsageError: Line magic function `%%time` not found.

```

In [66]: 1 # Summarize the results
2 print('Best : {}, using {}'.format(grid_result.best_score_,grid_result.best_params_))
3 means = grid_result.cv_results_['mean_test_score']
4 stds = grid_result.cv_results_['std_test_score']
5 params = grid_result.cv_results_['params']
6 for mean, stdev, param in zip(means, stds, params):
7     print('{} with: {}'.format(mean, stdev, param))

```

Best : -0.007342452276498079, using {'dropout_rate': 0.0, 'learning_rate': 0.001}
-0.007342452276498079,0.0019959824416733475 with: {'dropout_rate': 0.0, 'learning_rate': 0.001}
-0.008561298809945583,0.001991463301998141 with: {'dropout_rate': 0.0, 'learning_rate': 0.01}
-0.42658794578164816,0.48784841583609684 with: {'dropout_rate': 0.0, 'learning_rate': 0.1}
-0.010010520182549953,0.0011587275738232285 with: {'dropout_rate': 0.1, 'learning_rate': 0.001}
-0.010904456581920385,0.0023298427190362866 with: {'dropout_rate': 0.1, 'learning_rate': 0.01}
-0.5191736936569213,0.36482949841031104 with: {'dropout_rate': 0.1, 'learning_rate': 0.1}
-0.011447767540812493,0.0033028640505105367 with: {'dropout_rate': 0.2, 'learning_rate': 0.001}
-0.017722824588418007,0.004725459057919966 with: {'dropout_rate': 0.2, 'learning_rate': 0.01}
-0.7296398758888245,0.28865086924640976 with: {'dropout_rate': 0.2, 'learning_rate': 0.1}

```

In [67]: 1 # Defining the model
2 #get best value for neuron by hyperparameter tuning
3 %%time
4 def create_model_three_hidden_layers(neuron1,neuron2,neuron3,neuron4):
5     model = Sequential()
6     model.add(Dense(neuron1,input_dim = 6,kernel_initializer = 'uniform',activation = 'relu'))
7     model.add(Dense(neuron2,input_dim = neuron1,kernel_initializer = 'uniform',activation = 'relu'))
8     model.add(Dense(neuron3,input_dim = neuron2,kernel_initializer = 'uniform',activation = 'relu'))
9     model.add(Dense(neuron4,input_dim = neuron3,kernel_initializer = 'uniform',activation = 'relu'))
10    model.add(Dense(1))
11
12    adam = Adam(lr = 0.001)
13    model.compile(loss = 'mse',optimizer = adam,metrics = ['mse', 'mae', 'mape'])
14    return model
15
16 # Create the model
17
18 model = KerasRegressor(build_fn = create_model_three_hidden_layers,verbose = 0,batch_size = 70,epochs = 300)
19
20 # Define the grid search parameters
21
22 neuron1 = [8,16,32]
23 neuron2 = [32,64,128]
24 neuron3 = [32,64,128]
25 neuron4 = [32,64,128]
26
27 # Make a dictionary of the grid search parameters
28
29 param_grids = dict(neuron1 = neuron1,neuron2 = neuron2, neuron3 = neuron3, neuron4 = neuron4)
30
31 # Build and fit the GridSearchCV
32
33 grid = GridSearchCV(estimator = model,param_grid = param_grids,cv = KFold(),verbose = 0)
34 grid_result = grid.fit(x_train,y_train)

```

UsageError: Line magic function `%%time` not found.

```

In [68]: 1 # Summarize the results
2 print('Best : {}, using {}'.format(grid_result.best_score_,grid_result.best_params_))
3 means = grid_result.cv_results_['mean_test_score']
4 stds = grid_result.cv_results_['std_test_score']
5 params = grid_result.cv_results_['params']
6 for mean, stdev, param in zip(means, stds, params):
7     print('{} with: {}'.format(mean, stdev, param))

```

Best : -0.007342452276498079, using {'dropout_rate': 0.0, 'learning_rate': 0.001}

-0.007342452276498079,0.0019959824416733475 with: {'dropout_rate': 0.0, 'learning_rate': 0.001}

-0.008561298809945583,0.001991463301998141 with: {'dropout_rate': 0.0, 'learning_rate': 0.01}

-0.42658794578164816,0.48784841583609684 with: {'dropout_rate': 0.0, 'learning_rate': 0.1}

-0.010010520182549953,0.0011587275738232285 with: {'dropout_rate': 0.1, 'learning_rate': 0.001}

-0.010904456581920385,0.0023298427190362866 with: {'dropout_rate': 0.1, 'learning_rate': 0.01}

-0.5191736936569213,0.36482949841031104 with: {'dropout_rate': 0.1, 'learning_rate': 0.1}

-0.011447767540812493,0.0033028640505105367 with: {'dropout_rate': 0.2, 'learning_rate': 0.001}

-0.017722824588418007,0.004725459057919966 with: {'dropout_rate': 0.2, 'learning_rate': 0.01}

-0.7296398758888245,0.28865086924640976 with: {'dropout_rate': 0.2, 'learning_rate': 0.1}

```

In [69]: 1 #create a model with 3 hidden layers with best hyperparameters
2 def create_model_three_hidden_layers():
3     model = Sequential()
4     model.add(Dense(8, input_dim=6, kernel_initializer='uniform', activation='relu'))
5     model.add(Dense(128, kernel_initializer='uniform', activation='relu'))
6     model.add(Dense(64, kernel_initializer='uniform', activation='relu'))
7     model.add(Dense(128, kernel_initializer='uniform', activation='relu'))
8     model.add(Dense(1))
9
10    adam=Adam(lr=0.001)
11    model.compile(loss='mse', optimizer=adam, metrics=['mse', 'mae', 'mape'])
12    return model

```



```
In [70]: 1 %%time
2 epochs=300
3 batch_size=70
4
5 final_model=create_model_three_hidden_layers()
6
7 print("Here is the summary of our final model:")
8 final_model.summary()
9
10 with tf.device('/GPU:0'):
11     final_model.fit(x_train,y_train, verbose = 0,batch_size = batch_size,epochs = epochs, shuffle=True)
```

Here is the summary of our final model:

Model: "sequential_94"

| Layer (type) | Output Shape | Param # |
|-------------------|--------------|---------|
| dense_469 (Dense) | (None, 8) | 56 |
| dense_470 (Dense) | (None, 128) | 1152 |
| dense_471 (Dense) | (None, 64) | 8256 |
| dense_472 (Dense) | (None, 128) | 8320 |
| dense_473 (Dense) | (None, 1) | 129 |

```
=====
Total params: 17,913
Trainable params: 17,913
Non-trainable params: 0
```

Wall time: 12.4 s

```
In [77]: 1 loss, mae, mse, mape = final_model.evaluate(x_train, y_train)
2 print('\n', "Results for final model :", '\n', "Training Loss:", loss, '\n',
3       "Training Mean Absolute Error:" , mae, '\n', "Training Mean Squared Error:", mse)
```

36/36 [=====] - 0s 2ms/step - loss: 0.0059 - mse: 0.0059 - mae: 0.0552 - mape: 128.9680

Results for final model :
 Training Loss: 0.005904133897274733
 Training Mean Absolute Error: 0.005904133897274733
 Training Mean Squared Error: 0.055208414793014526

```
In [78]: 1 loss_t, mae_t, mse_t, mape_t = final_model.evaluate(x_test, y_test)
2 print('\n', "Results for final model :", '\n', "Test Loss:", loss_t, '\n',
3       "Test Mean Absolute Error:" , mae_t, '\n', "Test Mean Squared Error:", mse_t)
```

12/12 [=====] - 0s 3ms/step - loss: 0.0072 - mse: 0.0072 - mae: 0.0613 - mape: 101.6804

Results for final model :
 Test Loss: 0.007246457971632481
 Test Mean Absolute Error: 0.007246457971632481
 Test Mean Squared Error: 0.06132793426513672

Predicting values from Model using same dataset

```
In [85]: 1 from keras.wrappers.scikit_learn import KerasRegressor
```

```
In [87]: 1 # generating predictions for test data
2 y_predict_test = model1.predict(x_test)
3
4 # creating table with test price & predicted price for test
5 predictions_df = pd.DataFrame(x_test)
6 predictions_df['Actual'] = y_test
7 predictions_df['Predicted'] = y_predict_test
8 print(predictions_df.shape)
9 predictions_df.head(10)
```

12/12 [=====] - 0s 2ms/step
(376, 8)


Out[87]:

| | AFDP | GTEP | TIT | TAT | CDP | CO | Actual | Predicted |
|-------|-----------|-----------|-----------|-----------|-----------|-----------|-----------|---------------|
| 2317 | -0.066281 | -1.644945 | -2.146537 | 0.544795 | -1.667984 | 1.906030 | -1.931445 | 6.653878e-04 |
| 6485 | -0.843512 | -1.345066 | -1.598640 | 0.589392 | -1.416745 | 1.034617 | -1.474914 | 4.127419e-04 |
| 9091 | 0.119701 | -0.266125 | 0.178980 | 0.583021 | -0.033574 | -0.446970 | -0.004782 | -6.646042e-07 |
| 5443 | 0.368204 | 0.127002 | 0.501631 | 0.613602 | 0.076172 | -0.364359 | -0.028810 | -1.562580e-05 |
| 10027 | -1.553088 | -1.524370 | -1.793448 | 0.586844 | -1.649844 | 0.790223 | -1.484398 | 3.629955e-04 |
| 11794 | 0.426776 | 1.158241 | 0.866895 | -0.979151 | 1.101986 | -0.635490 | 0.800787 | -4.597998e-05 |
| 10773 | -1.538083 | -1.344586 | -1.556026 | 0.586844 | -1.370489 | 0.211813 | -1.402830 | 1.968982e-04 |
| 8915 | -0.049171 | -0.369920 | 0.227682 | 0.583021 | -0.084366 | -0.308247 | 0.098286 | 2.172192e-05 |
| 6588 | 0.166427 | -0.263248 | 0.185068 | 0.571554 | -0.091622 | -0.183720 | -0.035133 | 2.059035e-05 |
| 1384 | 0.707132 | -0.252461 | 0.075489 | 0.594489 | -0.251254 | -0.339798 | -0.095203 | 3.739018e-05 |

Visualizing the Relationship between the Actual and Predicted Values Model Validation

```
In [88]: 1 plt.figure(figsize=(12,8))
2 plt.xlabel("Actual Values")
3 plt.ylabel("Predicted values")
4 plt.title("The Scatterplot of Relationship between Actual Values and Predictions")
5 plt.scatter(predictions_df['Actual'], predictions_df['Predicted'])
```

Out[88]: <matplotlib.collections.PathCollection at 0x1417f9b7d00>



```
In [ ]: 1
```