

## Assignment-15-Random Forest (Fraud Data)

Use Random Forest to prepare a model on fraud data treating those who have taxable\_income <= 30000 as "Risky" and others are "Good"

### Import Libraries

```
In [1]: 1 import pandas as pd
2 import seaborn as sns
3 from matplotlib import pyplot as plt
4 import warnings
5 warnings.filterwarnings("ignore")
6
7 from sklearn.model_selection import train_test_split
8 from sklearn.metrics import confusion_matrix, classification_report, accuracy_score, roc_curve, roc_auc_score
```

### Import Data

```
In [2]: 1 data_fraud=pd.read_csv('Fraud_check.csv')
2 data_fraud
```

```
Out[2]:
```

	Undergrad	Marital.Status	Taxable.Income	City.Population	Work.Experience	Urban
0	NO	Single	68833	50047	10	YES
1	YES	Divorced	33700	134075	18	YES
2	NO	Married	36925	160205	30	YES
3	YES	Single	50190	193264	15	YES
4	NO	Married	81002	27533	28	NO
...	...	...	...	...	...	...
595	YES	Divorced	76340	39492	7	YES
596	YES	Divorced	69967	55369	2	YES
597	NO	Divorced	47334	154058	0	YES
598	YES	Married	98592	180083	17	NO
599	NO	Divorced	96519	158137	16	NO

600 rows × 6 columns

### Data Understanding

**Undergrad** : person is under graduated or not

**Taxable.Income** : Taxable income is the amount of how much tax an individual owes to the government

**Work Experience** : Work experience of an individual person


**Urban** : Whether that person belongs to urban area or not

```
In [3]: 1 data_fraud.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 600 entries, 0 to 599
Data columns (total 6 columns):
#   Column                Non-Null Count  Dtype
---  ---
0   Undergrad              600 non-null   object
1   Marital.Status         600 non-null   object
2   Taxable.Income         600 non-null   int64
3   City.Population        600 non-null   int64
4   Work.Experience        600 non-null   int64
5   Urban                  600 non-null   object
dtypes: int64(3), object(3)
memory usage: 28.2+ KB
```

```
In [4]: 1 fig,axes=plt.subplots(1,2)
2 plt.figure(figsize=[15,15])
3
4 fig.set_figheight(6)
5 fig.set_figwidth(15)
6
7 sns.kdeplot(x="City.Population", data=data_fraud,ax=axes[0])
8
9 sns.kdeplot(x="Work.Experience", data=data_fraud,ax=axes[1])
```


Out[4]: <AxesSubplot:xlabel='Work.Experience', ylabel='Density'>



<Figure size 1080x1080 with 0 Axes>

```
In [5]: 1 fig2,axes2=plt.subplots(1,2)
2 plt.figure(figsize=[5,15])
3
4 fig2.set_figheight(6)
5 fig2.set_figwidth(15)
6
7 sns.distplot(data_fraud["City.Population"],ax=axes2[0])
8
9 sns.distplot(data_fraud["Work.Experience"],ax=axes2[1])
```


Out[5]: <AxesSubplot:xlabel='Work.Experience', ylabel='Density'>



<Figure size 360x1080 with 0 Axes>

```
In [6]: 1 fig2,axes2=plt.subplots(1,2)
2 plt.figure(figsize=[5,15])
3
4 fig2.set_figheight(6)
5 fig2.set_figwidth(15)
6
7 sns.distplot(data_fraud["City.Population"],ax=axes2[0])
8
9 sns.distplot(data_fraud["Work.Experience"],ax=axes2[1])
```

Out[6]: <AxesSubplot:xlabel='Work.Experience', ylabel='Density'>



<Figure size 360x1080 with 0 Axes>

## Data Preparation

```
In [7]: 1 Income_List=[]
2 for i in data_fraud["Taxable.Income"]:
3     if i > 30000:
4         Income_List.append(1)
5     else:
6         Income_List.append(0)
```

```
In [8]: 1 Undergrad_List=[]
2 for i in data_fraud["Undergrad"]:
3     if i=="YES":
4         Undergrad_List.append(1)
5     else:
6         Undergrad_List.append(0)
```

```
In [9]: 1 Urban_List=[]
2 for i in data_fraud["Urban"]:
3     if i=="YES":
4         Urban_List.append(1)
5     else:
6         Urban_List.append(0)
```

```
In [10]: 1 data_fraud["Taxable.Income"]=Income_List
2 data_fraud.Undergrad=Undergrad_List
3 data_fraud.Urban=Urban_List
```

```
In [11]: 1 from sklearn.preprocessing import LabelEncoder
2 le = LabelEncoder()
```

```
In [12]: 1 data_fraud_copy=pd.read_csv('Fraud_check.csv')
```

```
In [13]: 1 data_fraud_copy["Marital.Status"]=le.fit_transform(data_fraud_copy["Marital.Status"])
```

```
In [14]: 1 data_fraud["Marital.Status"]=data_fraud_copy["Marital.Status"]
```

In [15]: 1 data\_fraud.head()

Out[15]:

	Undergrad	Marital.Status	Taxable.Income	City.Population	Work.Experience	Urban
0	0	2	1	50047	10	1
1	1	0	1	134075	18	1
2	0	1	1	160205	30	1
3	1	2	1	193264	15	1
4	0	1	1	27533	28	0

## Model Building

In [16]: 1 X=data\_fraud.drop("Taxable.Income",axis=1)  
2 y=data\_fraud[["Taxable.Income"]]

In [17]: 1 X\_train,X\_test,y\_train,y\_test=train\_test\_split(X,y,test\_size=.2,random\_state=123)  
2 print(X\_train.shape,y\_train.shape)  
3 print(X\_test.shape,y\_test.shape)

(480, 5) (480, 1)  
(120, 5) (120, 1)

## Model Training

In [18]: 1 from sklearn.ensemble import RandomForestClassifier  
2 rf\_model=RandomForestClassifier(max\_depth=3,random\_state=123)  
3 rf\_model.fit(X\_train,y\_train)

Out[18]: RandomForestClassifier(max\_depth=3, random\_state=123)

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.  
On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

In [19]: 1 rf\_model.score(X\_test,y\_test)

Out[19]: 0.7833333333333333

## Model Optimization

### GridSearch CV

In [20]: 1 from sklearn.model\_selection import GridSearchCV  
2 grid\_search\_cv = GridSearchCV(estimator = rf\_model,  
3 param\_grid={'criterion':['gini','entropy'],  
4 'max\_depth':[2,3,4,5,6]},  
5 cv=5)  
6 grid\_search\_cv.fit(X,y)  
7 print(grid\_search\_cv.best\_params\_)  
8 print(grid\_search\_cv.best\_score\_)

{'criterion': 'gini', 'max\_depth': 2}  
0.7933333333333332

In [21]: 1 rf\_model\_1 = RandomForestClassifier(criterion='gini',max\_depth=2,random\_state=123)  
2 rf\_model\_1.fit(X\_train,y\_train)

Out[21]: RandomForestClassifier(max\_depth=2, random\_state=123)

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.  
On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

## Evaluating rf\_Model\_1

In [22]: 1 rf\_model\_1.score(X\_test,y\_test)

Out[22]: 0.7833333333333333

```
In [23]: 1 y_pred_test = rf_model.predict(X_test)
```

```
In [24]: 1 print(confusion_matrix(y_test,y_pred_test))
```

```
[[ 0 26]
 [ 0 94]]
```

```
In [25]: 1 print(classification_report(y_test,y_pred_test))
```


	precision	recall	f1-score	support
0	0.00	0.00	0.00	26
1	0.78	1.00	0.88	94
accuracy			0.78	120
macro avg	0.39	0.50	0.44	120
weighted avg	0.61	0.78	0.69	120

```
In [26]: 1 accuracy_score(y_test,y_pred_test)
```

```
Out[26]: 0.7833333333333333
```

```
In [27]: 1 fpr, tpr, thresholds = roc_curve(y_train,rf_model.predict_proba (X_train)[:,1])
2
3 auc = roc_auc_score(y_train,rf_model.predict_proba (X_train)[:,1])
4 print(auc)
5
6 import matplotlib.pyplot as plt
7 plt.plot(fpr, tpr, color='red', label='Random Forest model ( area = %0.2f)%auc)
8 plt.plot([0, 1], [0, 1], 'k--')
9 plt.xlabel('False Positive Rate or [1 - True Negative Rate ]')
10 plt.ylabel('True Positive Rate')
11 plt.show()
```

```
0.7679372796238915
```




Eventhough the model has accuracy score of 0.78 the stability and specificity of the model is very low(determined from the values of precision and recall)

## Data Optimization

```
In [28]: 1 plt.figure(figsize=(8,5))
2 plt.pie(x=y_train.value_counts(),labels=['Good','Risky'],explode=[0.1,0.05],
3         autopct='%0.1f',colors=['darkseagreen','lightcoral'],shadow=True,textprops = {"fontsize":15})
4 plt.show
```

```
Out[28]: <function matplotlib.pyplot.show(close=None, block=None)>
```



### The Data is imbalance

```
In [29]: 1 from imblearn.over_sampling import SMOTE
2 balanced = SMOTE()
3
4 X_balanced , y_balanced = balanced.fit_resample(X,y)
```

```
In [30]: 1 Optimized_Data = X_balanced.copy()
2 Optimized_Data['y']=y_balanced
```

```
1 Optimized_Data.head()
```

```
In [31]: 1 X_train_Opt,X_test_Opt,y_train_Opt,y_test_Opt=train_test_split(X_balanced,y_balanced, test_size=0.2,random_state=123)
2
3 print(X_train_Opt.shape)
4 print(y_train_Opt.shape)
```

```
(761, 5)
(761, 1)
```

```
In [32]: 1 RF_Model_Opt=RandomForestClassifier(criterion='gini',max_depth=3,random_state=123)
```

```
In [33]: 1 RF_Model_Opt.fit(X_train_Opt,y_train_Opt)
```

```
Out[33]: RandomForestClassifier(max_depth=3, random_state=123)
```

**In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.  
On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.**

```
In [34]: 1 Optimized_Data['y_predicted']=RF_Model_Opt.predict(X_balanced)
```

```
In [35]: 1 accuracy_score(y_test_Opt,RF_Model_Opt.predict(X_test_Opt))
```

```
Out[35]: 0.6492146596858639
```

```
In [36]: 1 confusion_matrix(y_test_Opt,RF_Model_Opt.predict(X_test_Opt))
```

```
Out[36]: array([[66, 18],
               [49, 58]], dtype=int64)
```

## Test Data

In [37]: 1 `print(classification_report(y_test_Opt,RF_Model_Opt.predict(X_test_Opt)))`

```

              precision    recall  f1-score   support


     0       0.57         0.79         0.66         84
     1       0.76         0.54         0.63        107

 accuracy          0.65         0.65         0.65        191
 macro avg         0.67         0.66         0.65        191
 weighted avg      0.68         0.65         0.65        191

```

In [38]: 1 `fpr2, tpr2, thresholds2 = roc_curve(y_train_Opt,RF_Model_Opt.predict_proba (X_train_Opt)[: ,1])`  
 2  
 3 `auc2 = roc_auc_score(y_train_Opt,RF_Model_Opt.predict_proba (X_train_Opt)[: ,1])`  
 4 `print(auc2)`  
 5  
 6 `import matplotlib.pyplot as plt`  
 7 `plt.plot(fpr2, tpr2, color='red', label='Random Forest Model ( area = %0.2f)'%auc2)`  
 8 `plt.plot([0, 1], [0, 1], 'k--')`  
 9 `plt.xlabel('False Positive Rate or [1 - True Negative Rate]')`  
 10 `plt.ylabel('True Positive Rate')`  
 11 `plt.show()`

0.7609438084176761



The Model has precision and recall

## Model Deployment

In [39]: 1 `from pickle import dump,load`  
 2 `dump(RF_Model_Opt,open('Intel_On_frauddata_RF.pkl','wb'))`

In [40]: 1 `Loaded_Int=load(open('Intel_On_frauddata_RF.pkl','rb'))`  
 2 `Loaded_Int.predict(X_test_Opt.head())`

Out[40]: array([1, 1, 1, 1, 0], dtype=int64)

In [ ]: 1